

ECE3849
D-Term 2021

Real Time Embedded Systems

Module 1 Part 2

Module 1 Part 2 Overview

- Interrupts with Preemption Summary
- Rate-Monotonic Scheduling Theory
 - RMS Graphical Model
 - Response-Time Analysis
 - Processor Utilization
 - Utilization Upper Bound

Preemptive Strategy Summary

- Canonical System assumptions / rules
 1. All events are periodic.
 2. The relative deadline = period.
 3. The events are not phase aligned and can happen at any time relative to each other.
- Preemptive strategy with interrupts.
 - Events with the shortest period are prioritized the highest.
 - Higher priority tasks preempt/interrupt lower priority tasks.
 - The strategy provides the **lowest possible latencies** for the **highest priority** events.
 - Lower priority events still suffer from starvation and are at risk for missing their deadlines.
- Adds complexity requiring an interrupt controller to trigger events.

Preemptive Strategy Summary

- Event prioritization is important.
 - If all events have the same priority, the design behaves like priority polling.
 - Interrupts without priority have added complexity and no performance advantage.
- Prioritization strategy: shortest period = highest priority.
 - When interrupts fully enabled, highest priority event latency is very small.
 - Calculating response time of lower priority interrupts becomes difficult.
- Due to shared data and resources, we may need to disable interrupts.
 - When interrupts are disabled for the entire ISR, the design performs the same as priority polling.
 - The latency of all ISRs is affected when interrupts are disabled.
 - Latency of all events increases by the maximum time that interrupts are disabled.
 - It is very important to keep the time interrupts are disabled as short as possible.

Even small changes make a big difference.

```
165 //loop for testing the effect of disabling interrupts
166 while (true) {
167     IntMasterDisable();
168 //     delay_us(100);
169     count_unloaded++;
170     count_loaded--;
171     IntMasterEnable();
172
173     IntMasterDisable();
174 //     delay_us(200);
175     IntMasterEnable();
176 }
177 #endif
```

Original foreground while loop was empty.

Now we disable interrupts, for just two operations.

- It more than doubled the latency of the highest priority event, 600 nsec >> 266 nsec.

Original Latency without disabling ISRs

New latency

Expression	Type	Value	Value
(x)= event0_latency/120.0f	float	0.600000024	0.266666681
(x)= event1_latency/120.0f	float	2000.14172	2000.14172
(x)= event2_latency/120.0f	float	3001.21655	3001.25
(x)= event0_response_time/120.0f	float	2001.20837	2000.84998
(x)= event1_response_time/120.0f	float	3002.2251	3001.95825
(x)= event2_response_time/120.0f	float	12005.3838	12005.1504
(x)= event0_missed_deadlines	unsigned int	0	0
(x)= event1_missed_deadlines	unsigned int	0	0
(x)= event2_missed_deadlines	unsigned int	0	0

Better Analysis Methods

- When we introduced interrupts, response time of lower priority events becomes difficult to calculate.
 - Often they are preempted multiple times?
 - What if conditions change and event 1 can be preempted multiple times too.

Event	Period	Execution Time	Latency	Response Time (Latency + ?)	Relative Deadline (Period)	Schedulable ? YES = response < deadline
event0	6 ms	2 ms	0 ms	2 ms	6 ms	Yes
event1	8 ms	1 ms	2 ms	3 ms Tricky to calc	8 ms	Yes
event2	12 ms	6 ms	3 ms	12 ms Tricky to calc	12 ms	Maybe, marginal

- We need a more robust method of calculating response time to guarantee we have thought of the worst case scenario.

Rate-Monotonic Scheduling

Theory

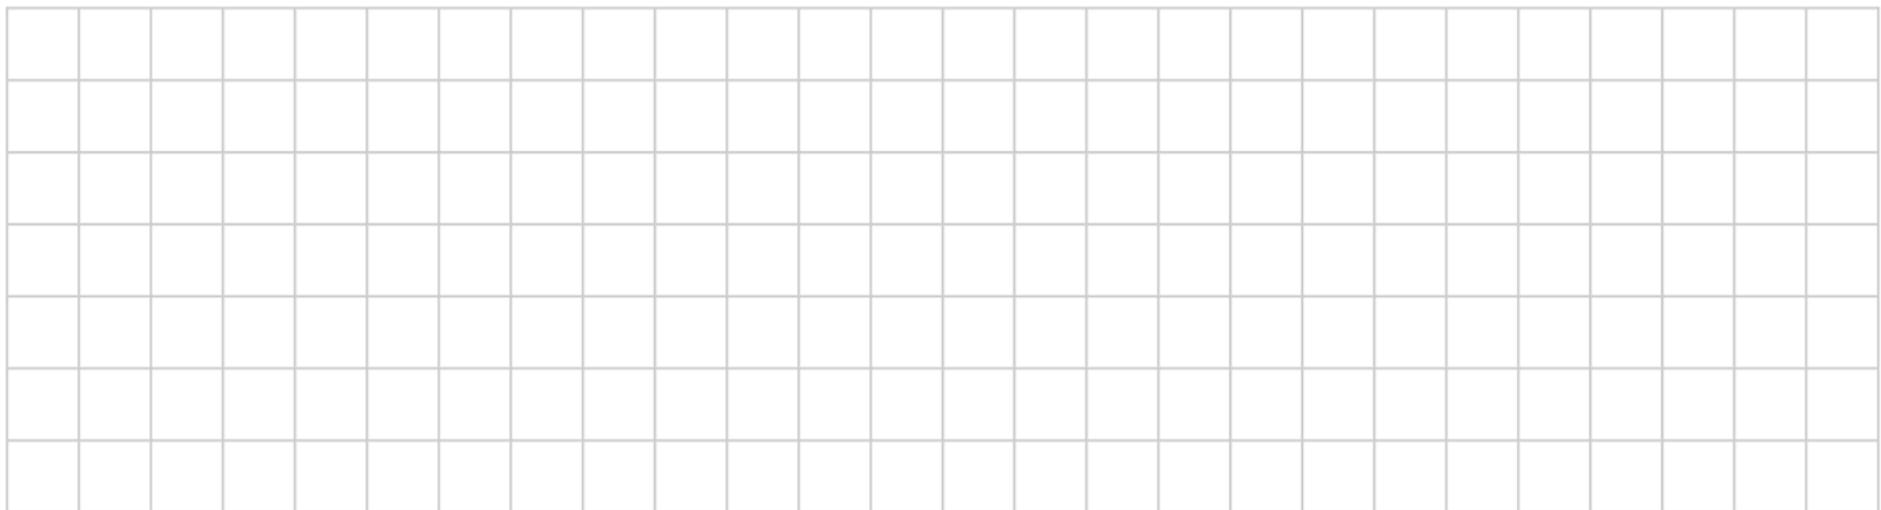
- Rate-Monotonic Scheduling (RMS) Theory provides a more reliable way to verify the scheduling.
- The system must satisfy the following conditions.
 - All real-time tasks (events) are periodic with fixed periods.
 - Relative deadline = event period.
 - Fixed priorities are assigned to events according to period.
 - Shortest period = higher priority.
 - Higher priority tasks preempt lower priority ones.
 - Execution time per event is fixed for each task.
 - This is accomplished by using the maximum value.
 - Task switching overhead is negligible.
 - Tasks do not synchronize with each other
 - No blocking or disabling interrupts.

RMS – Theorem

- Theorem: “The first deadline is the hardest”.
- If a task meets its first deadline when all tasks are released at the same time, then it will meet all future ones.
 - Worst case when all the tasks happen at once.
- Look at each task.
 - Does it meet its first deadline?
 - If yes, it is schedulable for all conditions.
- Lets use a graphical model.

Task	Period	Execution Time	Priority	Latency	Response Time	Schedulable?
task0	6 ms	2 ms				
task1	8 ms	1 ms				
task2	12 ms	6 ms				

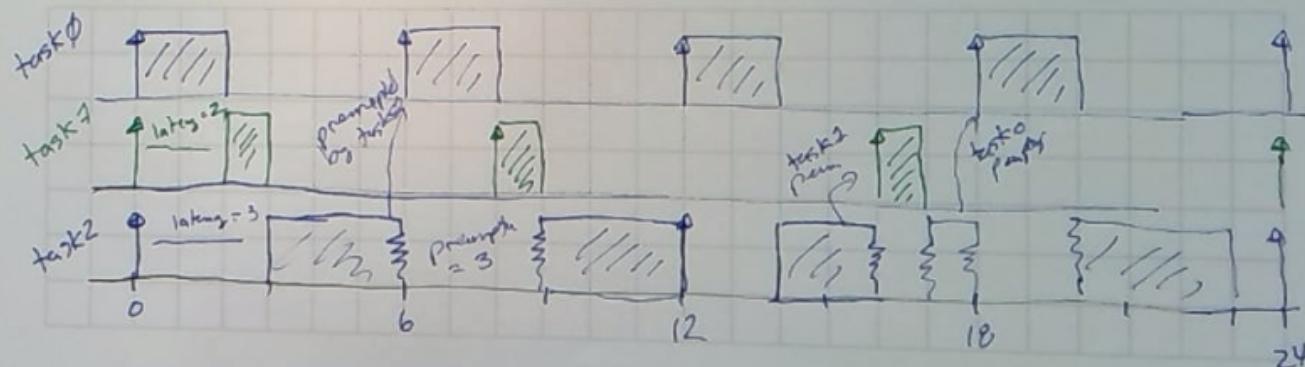
Unrolled schedule:



In-Class Worksheet

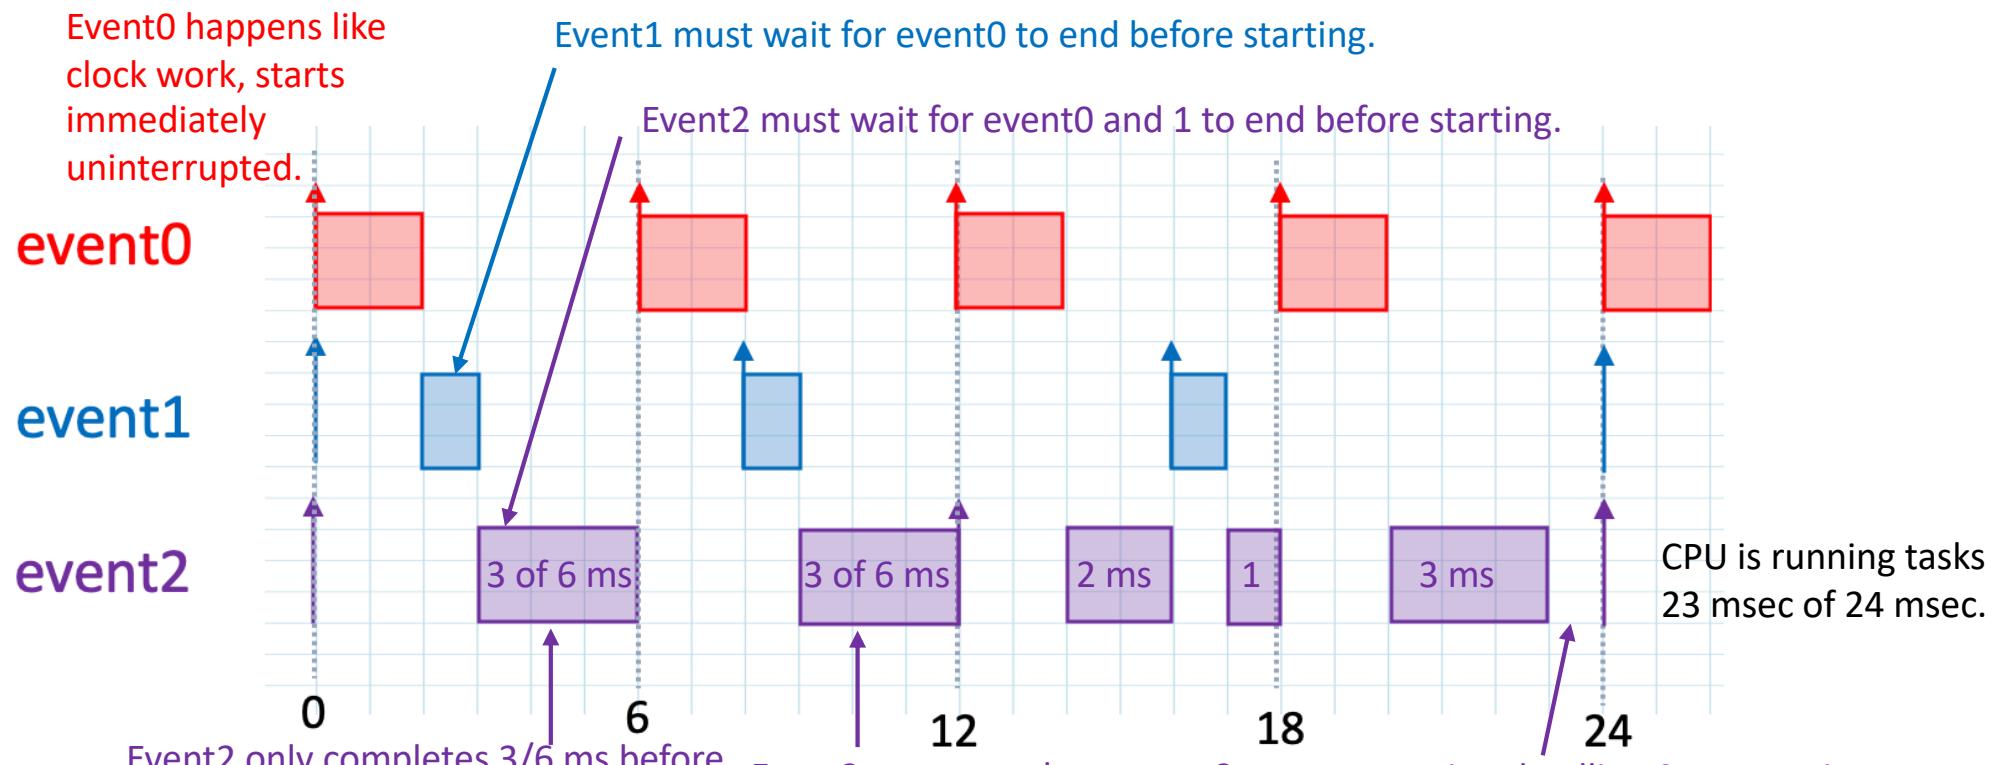
Task	Period	Execution Time	Priority	Latency	Response Time	Schedulable?
task0	6 ms	2 ms	high	0	$0+2=2$	$2 < 6$ Yes
task1	8 ms	1 ms	mid	2	$2+1=3$	$3 < 8$ Yes
task2	12 ms	6 ms	low	3	$3 \text{ latency} + 6 \text{ response}$ $+ 3 \text{ pre-empt} = 12 \text{ ms}$	$12 = 12$ margin

Unrolled schedule:



RMS Graphical Model

Event	Period	Execution Time	Priority	Latency	Response Time (Latency + ?)	Relative Deadline (Period)	Schedulable ? YES = response < deadline
event0	6 ms	2 ms	High	0 ms	2 ms	6 ms	Yes
event1	8 ms	1 ms	Mid	2 ms	3 ms	8 ms	Yes
event2	12 ms	6 ms	Low	3 ms	12 ms	12 ms	Yes, marginal



Response-Time Analysis

- The graphical method is difficult for non-round number and as the number of tasks grow.
- Given the execution time and period of each task, the response time can be calculated using iterative equation.

T_i = period of task i

E_i = execution time per event of task i

R_i = response time of task i

$i = 1$ for high priority task

$i = 2$ for mid priority task

$i = 3$ for low priority task

- Initialize the response time.

$$R_i^1 = \sum_{j=1}^i E_j$$

$$R_1^1 = E_1$$

$R_2^1 = E_1 + E_2$ (best case assumes task 1 only interrupts task 2 once)

$R_3^1 = E_1 + E_2 + E_3$ (best case assumes

task 1 and task 2 only interrupt task 3 once)

- Repeat this equation until it converges, where k is the number of iterations.

$$R_i^{k+1} = E_i + \sum_{j=1}^{i-1} \left\lceil \frac{R_i^k}{T_j} \right\rceil E_j$$

$$R_1^2 = E_1$$

$$R_2^2 = E_2 + \text{ceil}(R_1^2/T_1) * E_1$$

$$R_3^2 = E_3 + \text{ceil}(R_1^2/T_1) E_1 + \text{ceil}(R_1^2/T_2) * E_2$$

Adjustment for multiple interrupts

Response-Time Analysis: Matlab

- This function is implemented in Matlab.

MATLAB function implementing this formula (available on the course website):

```
function [R,U] = response_time(T,E)
R = zeros(size(T));
for i = 1:length(R)
    R(i) = sum(E(1:i));
    R_old = 0;
    j = 1:i-1;
    while abs(R(i) - R_old) > 10*eps(R(i))
        R_old = R(i);
        R(i) = E(i) + sum(ceil(R(i)./T(j)) .* E(j));
    end
end
U = sum(E./T); % CPU utilization
```

Loops through each task.

Initializes the response time.

$$R_i^1 = \sum_{j=1}^i E_j$$

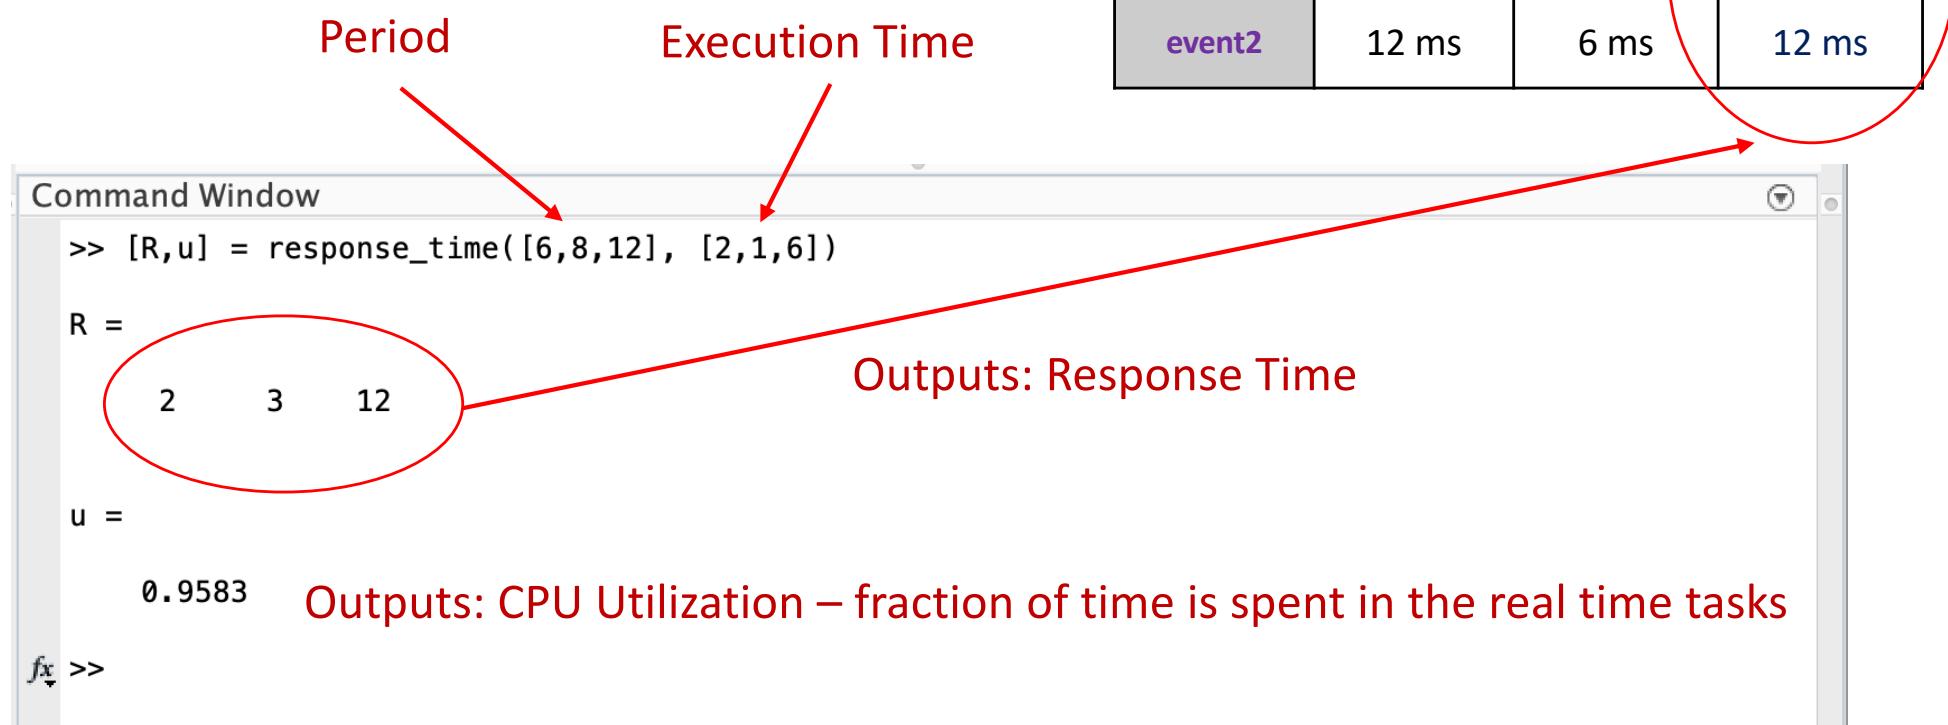
Loops until the difference between loops is very small.

Calculates response time.

$$R_i^{k+1} = E_i + \sum_{j=1}^{i-1} \left\lceil \frac{R_i^k}{T_j} \right\rceil E_j$$

Matlab Example Run

- Matlab file is named “response_time.m”.
- The file is in the search path or current directory in Matlab.



Theorem: Processor Utilization

- CPU utilization defines what fraction of the time is spent performing tasks.
 - It can be calculated by summing the fraction of time each task executes in its given period.

$$U = \sum_{i=1}^n \frac{E_i}{T_i}$$

where n is the number of task
E is the execution time.
T is period of task.

- CPU Utilization in previous example.

- $U = 2/6 + 1/8 + 6/12$
- $U = 23/24 = 0.958$

Event	Period (T)	Execution Time (E)	
event0	6 ms	2 ms	2 ms
event1	8 ms	1 ms	3 ms
event2	12 ms	6 ms	12 ms

Processor Utilization Upper Bound

- Given the RMS assumptions an upper bound for utilization is derived given a certain number of tasks.
 - Where n is the number of tasks.
 - If utilization is less than this, it can definitely be scheduled.
 - If utilization is more, that it may still be possible but further analysis is needed.
 - The last example had 3 tasks with a 0.958 utilization which is greater than the 0.7798 upper bound but it was schedulable.
 - In practice, we can not use 100% of the CPU for real time events. Generally it is good to have 10 – 20% margin.

$$U \leq n(2^{1/n} - 1)$$

n	$\max(U)$	$\max(U), \%$
1	1	100%
2	0.8284	82.84%
3	0.7798	77.98%
4	0.7568	75.68%
5	0.7435	74.35%
∞	$0.6931 = \ln(2)$	69.31%