

ECE3849
D-Term 2021

Real Time Embedded Systems

Module 1 Part 1

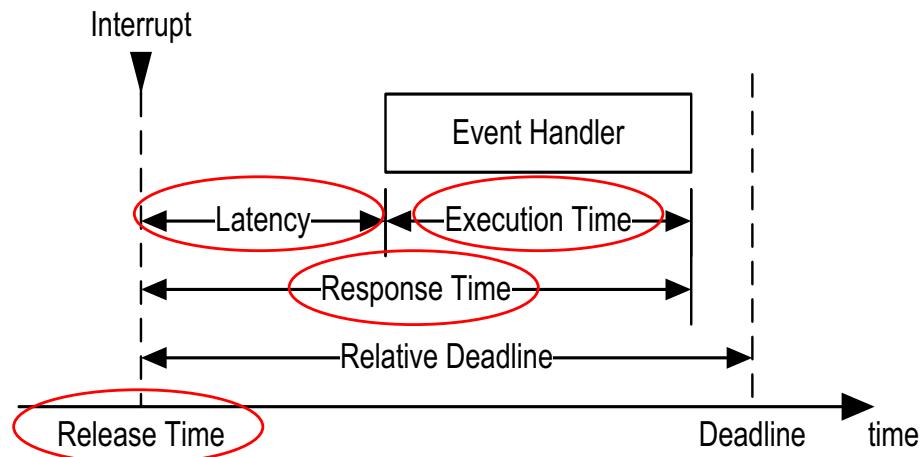
Module 1 Part 1 Overview

- Meeting and Setting deadlines
 - Latency
 - Execution Time
 - Response Time
- Scheduling Strategies
 - Round Robin Polling
 - Priority Polling
 - Interrupts with Preemption
- Rate-Monotonic Scheduling Theory

Events without Preemption

- Simplest case: No preemption

- The event handler / task runs from start to finish without interruption.
- Conditions this occurs under
 - Only one task running.
 - The task is highest priority.
- Event is serviced immediately.



- Release time
 - Hardware notifies the software that an external event has occurred.

- Latency
 - The time between when the event / interrupt happens and when the event handler / task starts to execute.
 - In this simple case depends on the microcontroller being used and the time for the software to context switch.

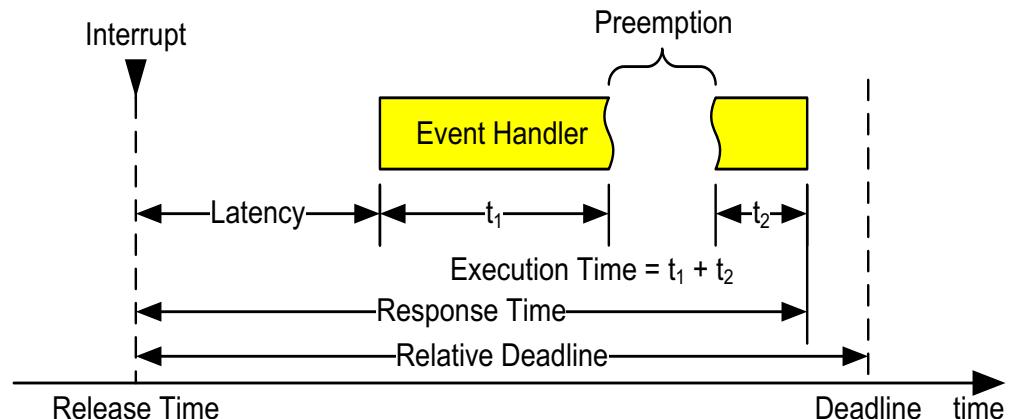
- Execution time depends on the event handler code and the hardware running the code.

- Fast hardware => faster execution time.
- Software should always optimize for the worst case.
- For scheduling hard deadlines, the average has little importance.

- Response time = Maximum Latency + Maximum Execution Time.

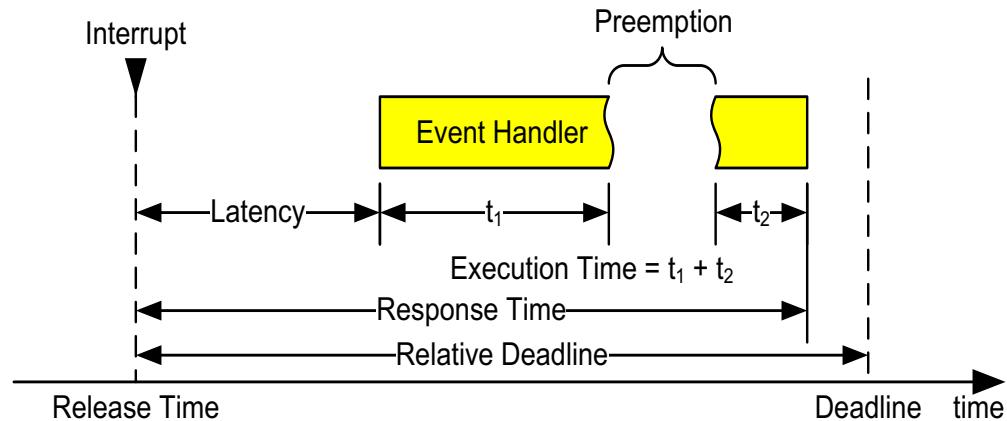
Events with Preemption

- There are multiple tasks
 - Each task has a priority.
 - Higher priority tasks can interrupt lower priority task.



- Release time
 - Same as before: Hardware notifies the software that an external event has occurred.
- Latency
 - Lower priority tasks may have to wait for higher priority tasks to complete before it starts.
 - Best case it starts immediately.
 - **Worst case** has to wait for other tasks to complete.
- Execution time now depends not only the event handler / task but also execution times of higher priority tasks.
 - Best case it executes without preemption, start to finish.
 - **Worst case** has to wait for all other higher priority task to complete.
- Response time = Maximum Latency + Maximum Execution time.

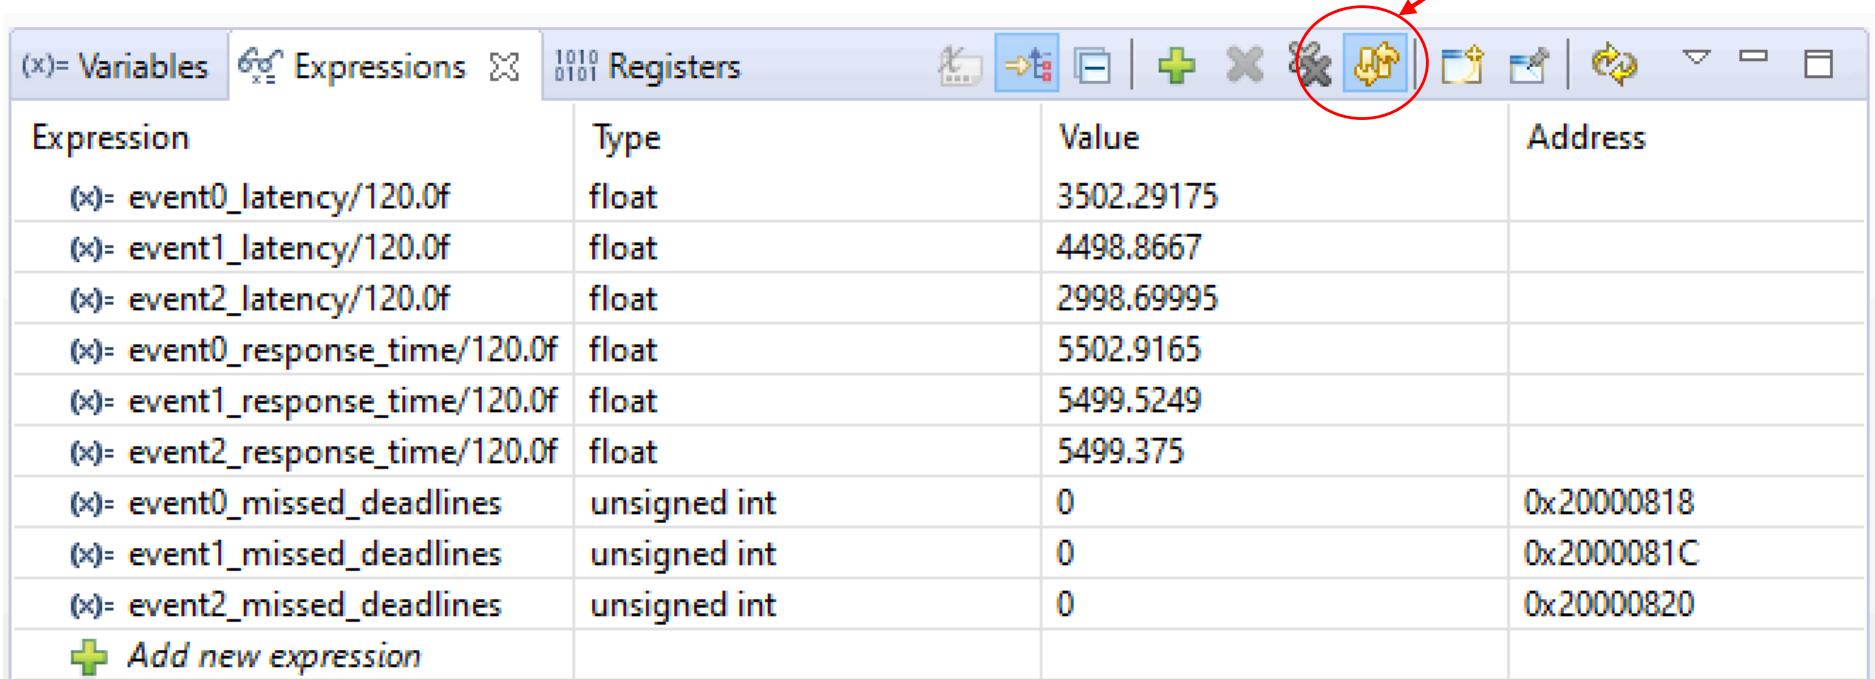
Meeting Deadlines



- The software must respond by the deadline.
 - Some deadlines are absolute, it must happen by 1:00 pm today.
 - Relative deadlines are defined relative to the release time.
 - When considering how to schedule a task, relative deadlines are important.
 - For periodic events the relative deadline equals the period of the event, as task must complete before the next event occurs.
- Relative deadline
 - Hard deadline – system fails if it misses any deadline . Some Examples?
 - Power cycle indicator (depends on the consequence)?
 - Thrusters on Mars rover when landing – come down too fast and lose the rove.
 - Soft deadline – systems tolerate if it misses some deadlines – Some Examples?
 - Non-critical sensor reading, home thermometer? Lab submissions – only because late deduction.
- A task is “schedulable” if the Maximum Response Time < Relative Deadline
- What might cause the response time to vary?
 - Performance throttling of hardware , slows down? Sequence of tasks change the average.
 - Conditional equations: if / else. If takes 1 msec, else 10 msec.

CCS Project: ece3849_int_latency

- Program: `ece3849_int_latency`
 - Generates three events: event0, event1 and event2
 - Each event has a
 - Period specified in usec
 - Execution time specified in usec
 - Measures the latency and execution time of each event in clock ticks.
 - The clock ticks are divided by 120 to convert to usec, because the clock is running at 120 MHz.
 - Increments a counter if any deadlines are missed.
- Results are viewed by adding expressions and using continuous refresh



The screenshot shows the CCS (Code Composer Studio) interface with the 'Expressions' tab selected. A red circle highlights the 'Continuous Refresh' button (a blue icon with a circular arrow) in the toolbar above the table. The table lists various expressions, their types, values, and addresses.

| Expression | Type | Value | Address |
|---|--------------|------------|------------|
| <code>(x)= event0_latency/120.0f</code> | float | 3502.29175 | |
| <code>(x)= event1_latency/120.0f</code> | float | 4498.8667 | |
| <code>(x)= event2_latency/120.0f</code> | float | 2998.69995 | |
| <code>(x)= event0_response_time/120.0f</code> | float | 5502.9165 | |
| <code>(x)= event1_response_time/120.0f</code> | float | 5499.5249 | |
| <code>(x)= event2_response_time/120.0f</code> | float | 5499.375 | |
| <code>(x)= event0_missed_deadlines</code> | unsigned int | 0 | 0x20000818 |
| <code>(x)= event1_missed_deadlines</code> | unsigned int | 0 | 0x2000081C |
| <code>(x)= event2_missed_deadlines</code> | unsigned int | 0 | 0x20000820 |
| + Add new expression | | | |

ece3849_int_latency

Defining Parameters and initializing values

```
33 // event and handler definitions
34 #define EVENT0_PERIOD          6007    // [us] event0 period
35 #define EVENT0_EXECUTION_TIME  2000    // [us] event0 handler execution time
36
37 #define EVENT1_PERIOD          8101    // [us] event1 period
38 #define EVENT1_EXECUTION_TIME  1000    // [us] event1 handler execution time
39
40 #define EVENT2_PERIOD          12301   // [us] event2 period
41 #define EVENT2_EXECUTION_TIME  2500    // [us] event2 handler execution time
42
43 // build options
44 // #define DISABLE_INTERRUPTS_IN_ISR // if defined, interrupts are disabled in the body of each ISR
45
46 // measured event latencies in clock cycles
47 uint32_t event0_latency = 0;
48 uint32_t event1_latency = 0;
49 uint32_t event2_latency = 0;
50
51 // measured event response time in clock cycles
52 uint32_t event0_response_time = 0;
53 uint32_t event1_response_time = 0;
54 uint32_t event2_response_time = 0;
55
56 // number of deadlines missed
57 uint32_t event0_missed_deadlines = 0;
58 uint32_t event1_missed_deadlines = 0;
59 uint32_t event2_missed_deadlines = 0;
60
```

- Sets period and execution time for each event.
- Small value added to cause variation in phase.

For each event initializes

- Latency
- Response time
- Number of missed deadlines

ece3849 int latency

TM4C1294NCPDT Timer Summery

- CORRECTION – Changed from original slides posted.
- There are two General Purpose Timers (GPTM Modules)
 - There are 8 GPTM timer modules.
 - Each GPTM timer has two sub timers - TimerA and B can be used individually in 16-bit mode or together for 32-bit counts.
 - They can count up or down.
 - They can be a one shot timer or a periodic timer that reloads itself when the time is reached.
 - See datasheet for more details.
- This example uses 3 GPTM timers using submodule TimerA.
 - Each event uses an interval timer with an interrupt output to trigger the event handler task.
- The the timer units are in number of CPU clocks.
 - The CPU is running at 120 MHz.
 - There are a 120 clocks in 1 usec.

```
61 // timer periods in clock cycles (expecting 120 MHz clock)
62 // 120 clock cycles in 1 usec
63 #define TIMER0_PERIOD (120 * EVENT0_PERIOD)
64 #define TIMER1_PERIOD (120 * EVENT1_PERIOD)
65 #define TIMER2_PERIOD (120 * EVENT2_PERIOD)
```

ece3849 int latency

Configuring the Timer Interrupts

```
// initialize general purpose timers 0-2 for periodic inter
SysCtlPeripheralEnable(SYSCTL_PERIPH_TIMER0);
TimerDisable(TIMER0_BASE, TIMER_BOTH);
TimerConfigure(TIMER0_BASE, TIMER_CFG_PERIODIC);
TimerLoadSet(TIMER0_BASE, TIMER_A, TIMER0_PERIOD - 1);
TimerIntEnable(TIMER0_BASE, TIMER_TIMA_TIMEOUT);
IntPrioritySet(INT_TIMER0A, 0); // 0 = highest priority, 31 = lowest
IntEnable(INT_TIMER0A);

SysCtlPeripheralEnable(SYSCTL_PERIPH_TIMER1);
TimerDisable(TIMER1_BASE, TIMER_BOTH);
TimerConfigure(TIMER1_BASE, TIMER_CFG_PERIODIC);
TimerLoadSet(TIMER1_BASE, TIMER_A, TIMER1_PERIOD - 1);
TimerIntEnable(TIMER1_BASE, TIMER_TIMA_TIMEOUT);
IntPrioritySet(INT_TIMER1A, 32); // 0 = highest priority, 31 = lowest
IntEnable(INT_TIMER1A);

SysCtlPeripheralEnable(SYSCTL_PERIPH_TIMER2);
TimerDisable(TIMER2_BASE, TIMER_BOTH);
TimerConfigure(TIMER2_BASE, TIMER_CFG_PERIODIC);
TimerLoadSet(TIMER2_BASE, TIMER_A, TIMER2_PERIOD - 1);
TimerIntEnable(TIMER2_BASE, TIMER_TIMA_TIMEOUT);
IntPrioritySet(INT_TIMER2A, 64); // 0 = highest priority, 31 = lowest
IntEnable(INT_TIMER2A);

IntMasterEnable(); // comment for polled scheduling or CI

TimerEnable(TIMER0_BASE, TIMER_A); // comment for CPU load
TimerEnable(TIMER1_BASE, TIMER_A);
TimerEnable(TIMER2_BASE, TIMER_A);
```

- Enables the Timer0 peripheral to be used.
- Disables operation during configuration.
- Counts down from the TIMERx_PERIOD
- Enables the timer interrupt to trigger when it reaches 0.
- Sets the priority of the interrupt.
- Enables the CPU to respond to the timer interrupt.

These function calls are documented in in the TivaWare Peripheral Driver Library User's Guide.

Enables all counters once configuration is complete

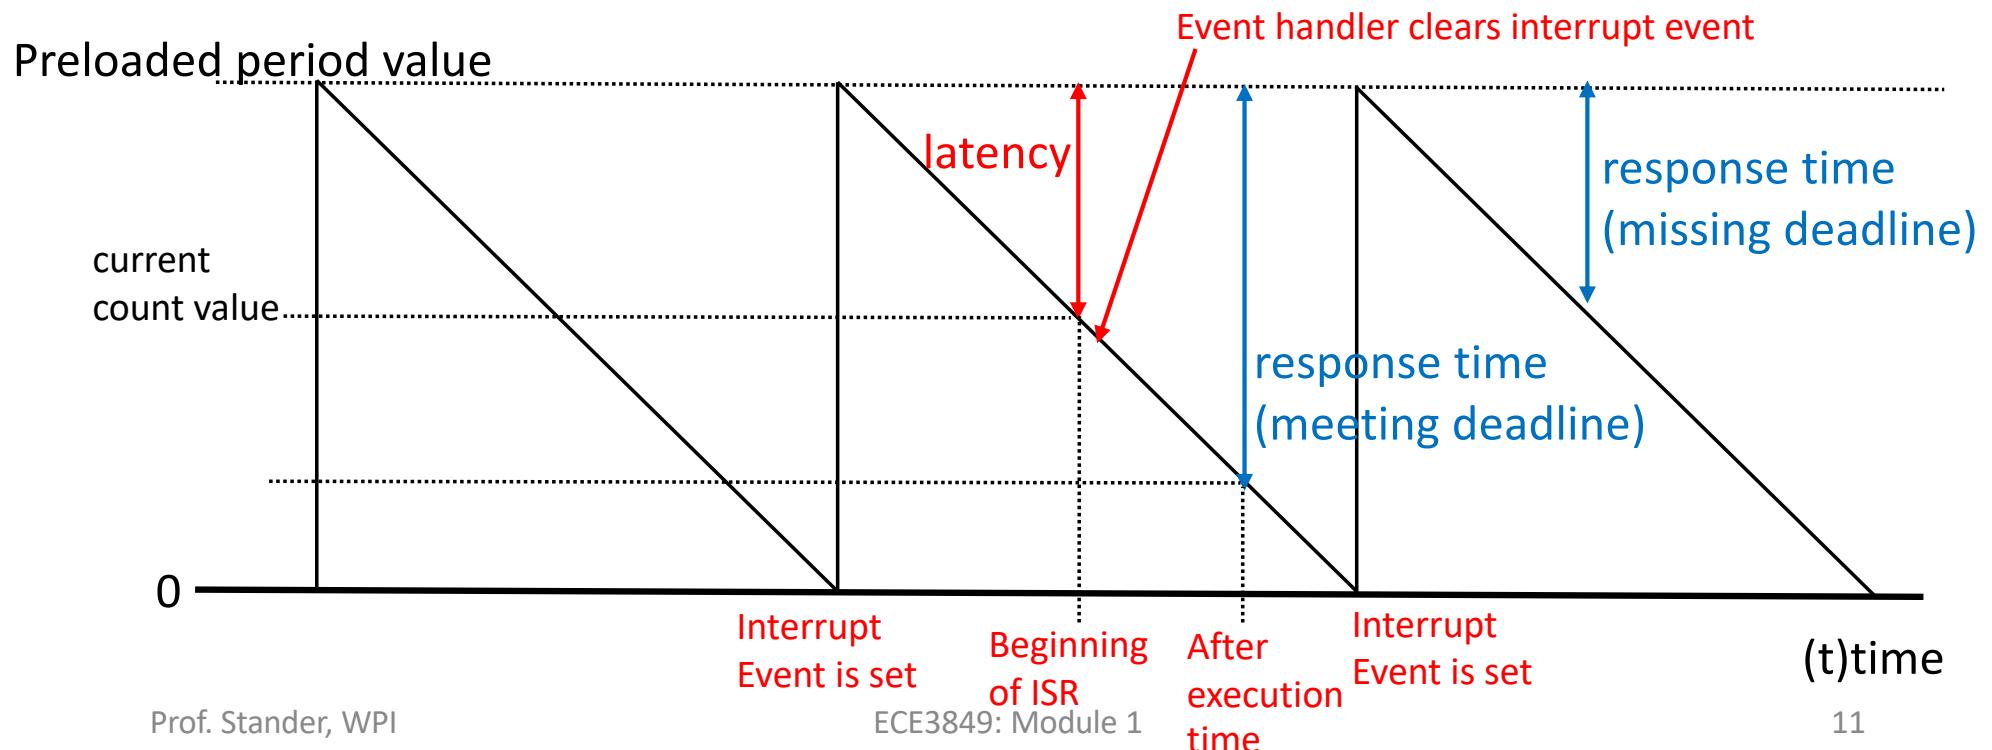
ece3849_int_latency

Infinite while loop

- There are several different scheduling modes we will explore.
 - Each will have different behaviors but all will have an infinite while loop.
- While in the loop, the timers control when events happen.
 - When an event happens an interrupt bit is set.
 - An event handler routine is called to service the event.
- The event handler does the following.
 - Reads the current time.
 - Clears the interrupt status register.
 - Calculates the latency.
 - Waits for the execution time.
 - Calculates the response time.

Calculating latency and response time.

- During configuration the preloaded period value is programmed.
- Timer decrements by 1 every clock.
- It counts down to zero then sets its interrupt status and reloads itself.
 - The release time is when the current count value = preloaded period value.
- Upon entering the event handler the interrupt event is cleared.
- Time_since_interrupt = timer_period – current_timer_count
 - Latency = preloaded period value – current timer value at the start of the interrupt.
 - Response time = preloaded period value – current timer value after the execution time delay.
 - If deadline is missed, the interrupt event is set.
 - We account for this by adding an additional preloaded period value.



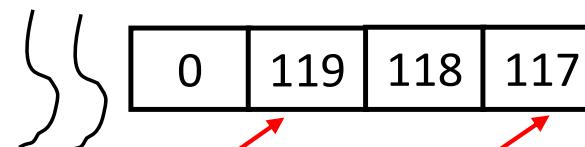
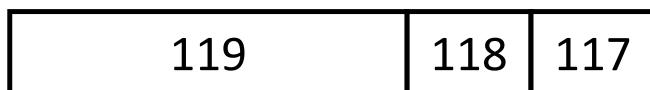
Timer Functional Calls Used

#1 TimerLoadSet(TIMERO_BASE, TIME_A, 119)

This examples sets timeout count 1usec = 120 clock counts

#2 TimerEnable(TIMERO_BASE, TIMER_A)

Enables counting down – 1 count per CPU clock edge



#3 Interrupt bit is set when timer reach 0 and then rolls back to pre-set value.

#4 TimerIntStatus(TIMERO_BASE, 1)

Reads the interrupt status register.



#5 In event handler gets the time
TimerValueGet(TIMERO_BASE, TIMER_A)

#6 Time in usec since last interrupt =
(119 – TimerValueGet output)/120

#7 TimerIntClear(TIMERO_BASE, TIMER_TIMA_TIMEOUT)
Clears Interrupt bit in event handler

ece3849_int_latency: Calculating Latency

- Each event handler is similar but accessing different hardware registers and updates different variables.

- Calculating Latency

- TIMER0_TAR_R is the hardware register containing the current time.
 - Latency is the difference between the period and the current time.
 - If it is the largest latency, then it updates the value.

```
193 void event0_handler(void)
194 {
195 #ifdef DISABLE_INTERRUPTS_IN_ISR
196     IntMasterDisable();
197 #endif
198     uint32_t t = TIMER0_PERIOD - TIMER0_TAR_R; // read
199     if (t > event0_latency) event0_latency = t; // meas
200     TimerIntClear(TIMER0_BASE, TIMER_TIMA_TIMEOUT); //
201
202     delay_us(EVENT0_EXECUTION_TIME); // handle event0
203 }
```

- Clears the interrupt, so it can be triggered again.
 - Sits in a loop for the execution time
- **TIMER0_TAR_R** is defined in tm4c1294ncpdt.h.
 - It defines register names that map to hardware register addresses.
- Hardware Registers defined in the **TM4C129NCPDT Microcontroller Datasheet**.

ece3849_int_latency

Calculating Response Time

- Handling a missed deadline
 - If the deadline was exceeded the counter rolls over and a new interrupt is generated.
 - The missed deadlines value is incremented
 - An extra period is added to the start count due to the counter rolling over.

- If the deadline was not missed, the original count is just the period.

```
203  
204 if (TimerIntStatus(TIMER0_BASE, 1) & TIMER_TIMA_TIMEOUT) {  
205     event0_missed_deadlines++;  
206     t = 2 * TIMER0_PERIOD; // timer overflowed since last interrupt  
207 }  
208 else t = TIMER0_PERIOD;  
209 t -= TimerValueGet(TIMER0_BASE, TIMER_A); // read Timer A value  
210 if (t > event0_response_time) event0_response_time = t; //
```

- The execution time is the period start count – the current time value
 - `TimerValueGet` function is equivalent in function accessing `TIMER0_TAR_R` directly but uses the API.
- If the response time is the maximum value the response time is updated.

ece3849 int latency

TM4C1294NCPDT Timer Summery

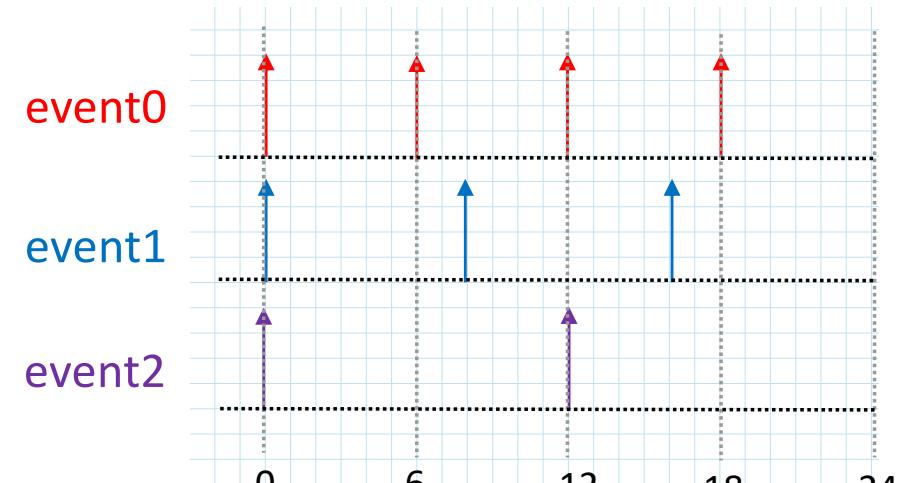
- CORRECTION !!!!!!
- There are two General Purpose Timers (GPTM Modules)
 - There are 8 GPTM timer modules.
 - Each GPTM timer has two sub timers - TimerA and B can be used individually in 16-bit mode or together for 32-bit counts.
 - They can count up or down.
 - They can be a one shot timer or a periodic timer that reloads itself when the time is reached.
 - See datasheet for more details.
- This example uses 3 GPTM timers using submodule TimerA.
 - Each event uses an interval timer with an interrupt output to trigger the event handler task.
- The the timer units are in number of CPU clocks.
 - The CPU is running at 120 MHz.
 - There are a 120 clocks in 1 usec.

```
61 // timer periods in clock cycles (expecting 120 MHz clock)
62 // 120 clock cycles in 1 usec
63 #define TIMER0_PERIOD (120 * EVENT0_PERIOD)
64 #define TIMER1_PERIOD (120 * EVENT1_PERIOD)
65 #define TIMER2_PERIOD (120 * EVENT2_PERIOD)
```

Canonical Real-Time Systems

- System Assumptions / Rules
 1. All Events are periodic.
 2. The relative deadline = period.
 - The event needs to finish before it can be called again.
 3. The events are not phase aligned and can happen at anytime relative to each other.
- Example: There are three events: event0, event1 and event2

| Event | Period | Execution Time |
|--------|--------|----------------|
| event0 | 6 ms | 2 ms |
| event1 | 8 ms | 1 ms |
| event2 | 12 ms | 2.5 ms |

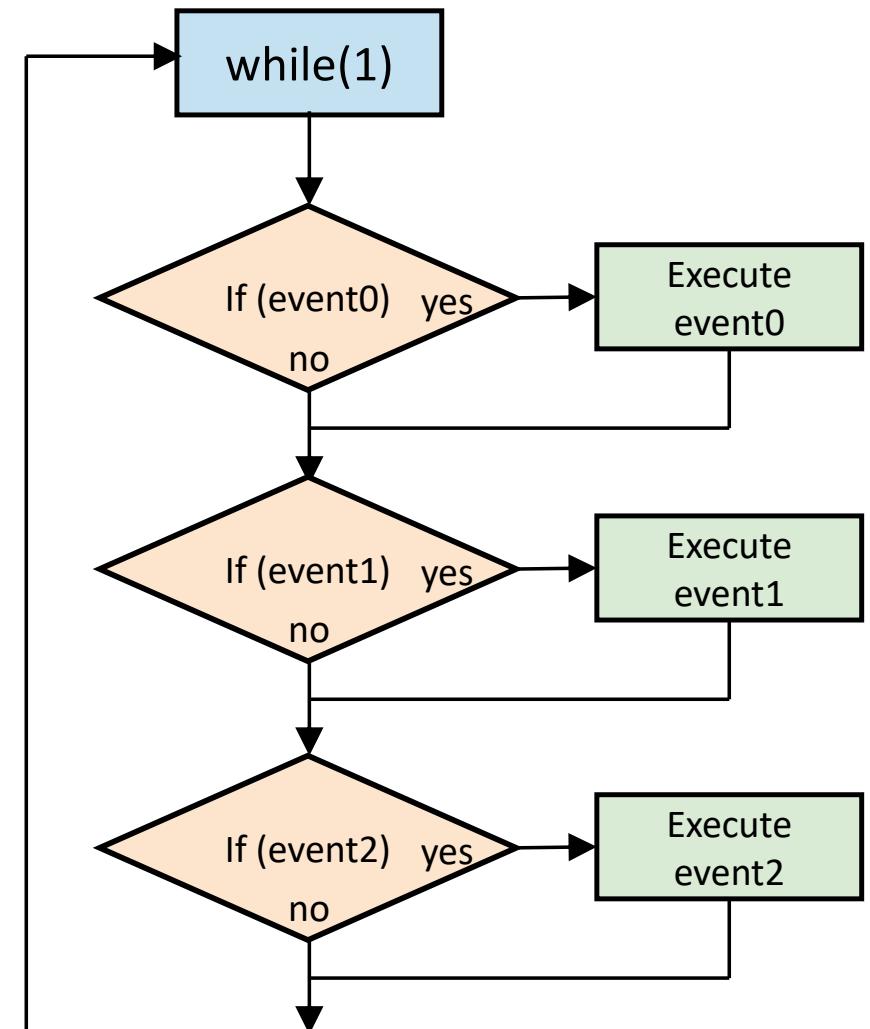


We have three events to service, any concerns?

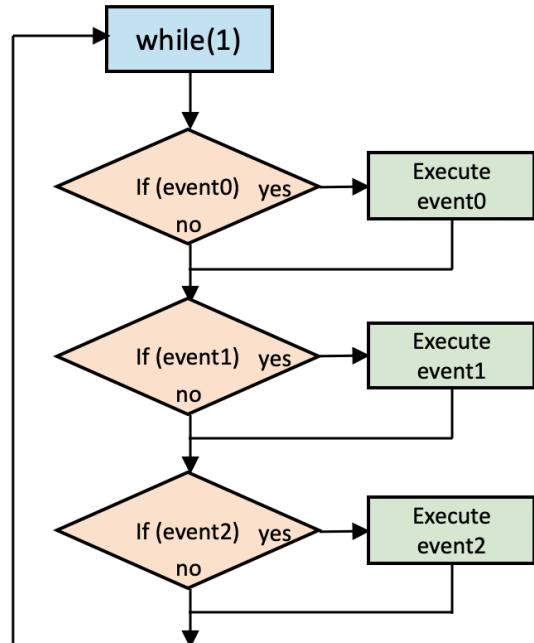
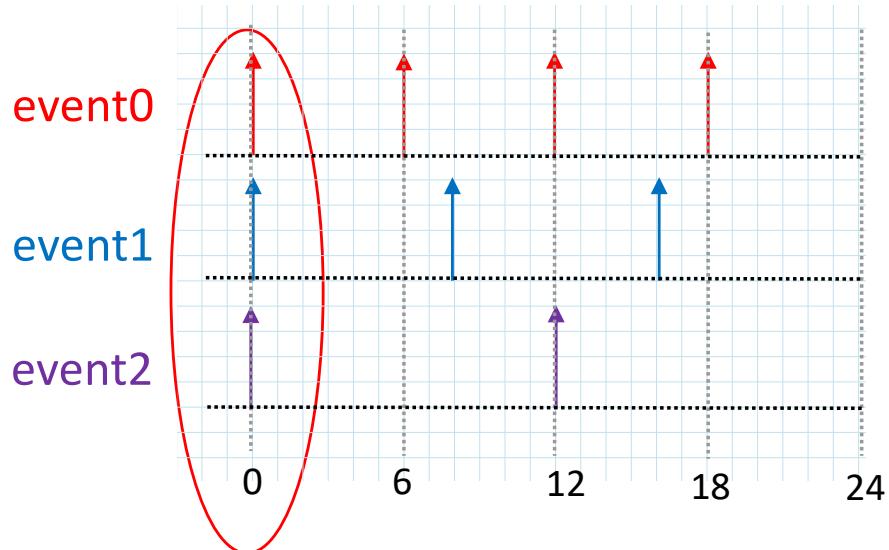
- How do we know if they make deadline?
- Are they preemption? Do we have it? We have not defined a scheduling algorithm.

Polling Loops Without Preemption

- Simplest way to run is a polling loop.
 - No priority => no preemption.
 - Each event is checked in sequence.
 - Each event is guaranteed to run. (starvation-free)
- Round-Robin Scheduling
 - The status of each event is checked (polled).
 - If an event occurred, it executes to completion.
 - Then repeats process for the next event.
 - When all events have been checked returns to the beginning



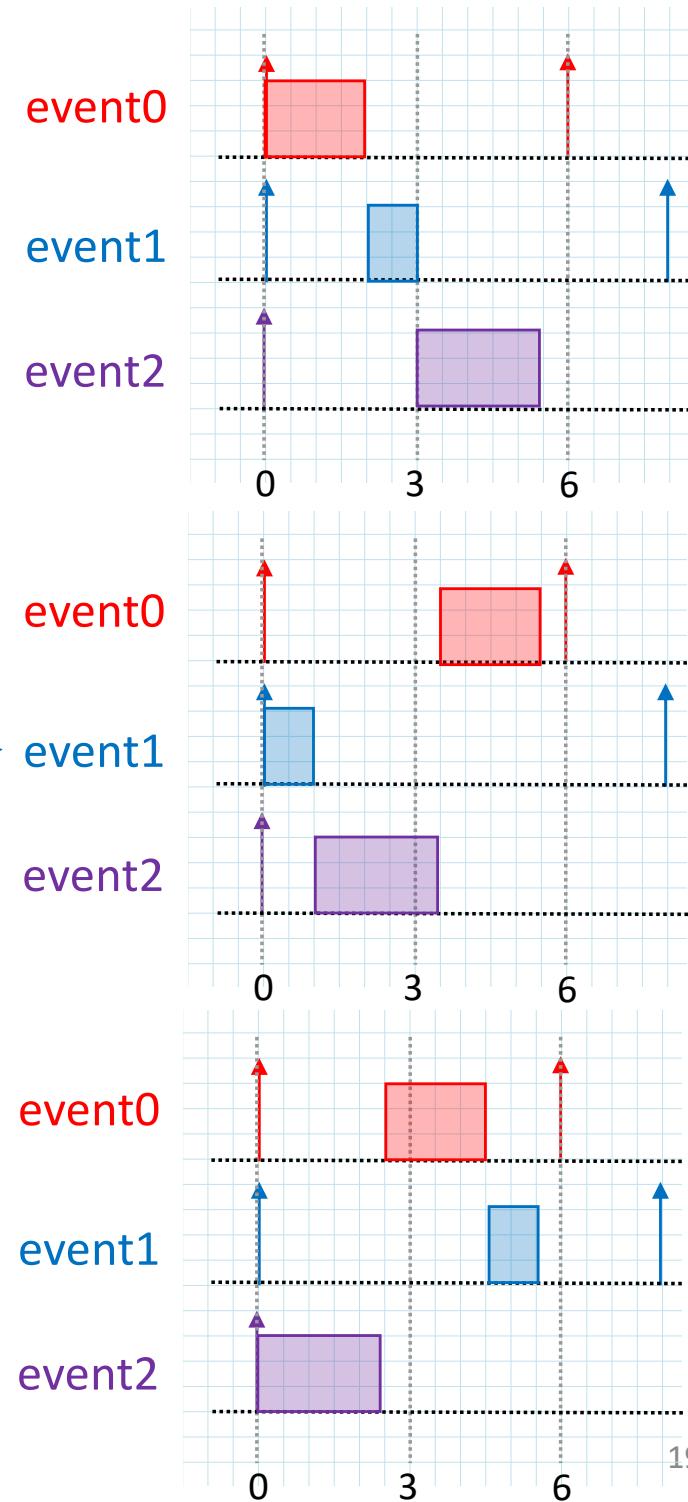
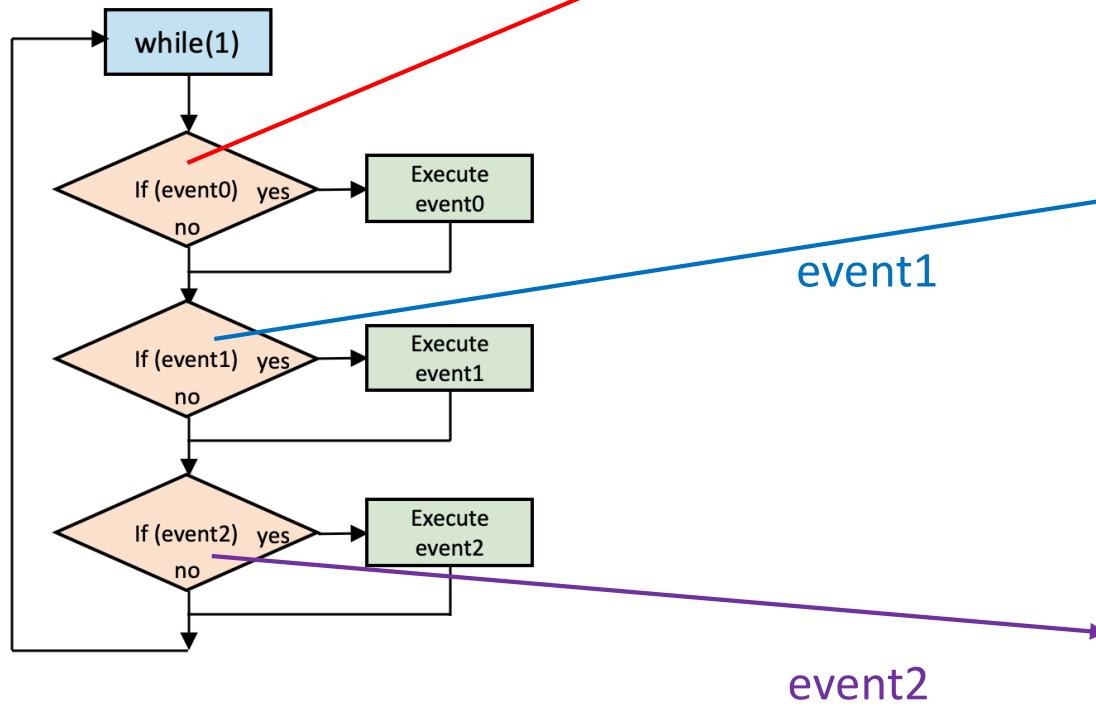
Round Robin Polling



- For the example,
 - All three events occur at relative time 0.
 - Polling loop has been running for a while.
- Which task get serviced first?
 - All equality important no priority?
 - We naturally want to 0 cause it is first, but loop can be any where in its execution.
 - We don't know.

Round Robin Scheduling

- The order of servicing tasks depends on where you are in the loop when they happen.



Are the tasks schedulable?

| Event | Period | Execution Time |
|--------|--------|----------------|
| event0 | 6 ms | 2 ms |
| event1 | 8 ms | 1 ms |
| event2 | 12 ms | 2.5 ms |

Remember, we care only about maximum value for these calculations.

```

void main (void) {
    <init>; // pseudo-code in <...>

    while (1) {
        if (<event0 occurred>) {
            <handle event0>;
        }
        if (<event1 occurred>) {
            <handle event1>;
        }
        if (<event2 occurred>) { texec0 = 2 ms
            <handle event2>;
        }
    }
}

```

$t_{exec0} = 2 \text{ ms}$

$t_{exec1} = 1 \text{ ms}$

$t_{exec2} = 2.5 \text{ ms}$

| Event | Period | Execution Time | Latency | Response Time | Relative Deadline | Schedulable ? |
|--------|--------|----------------|---|---|-----------------------------|---------------|
| event0 | 6 ms | 2 ms | 1 event 1 + event 2 = 1 + 2.5 = 3.5 | Latency+ execution max = 3.5 + 2m = 5.5m | 5.5 m < 6m | YES |
| event1 | 8 ms | 1 ms | $2 + 2.5 = 4.5\text{m}$ | $1 + 4.5 = 5.5$ | $5.5 \text{ m} < 8\text{m}$ | YES |
| event2 | 12 ms | 2.5 ms | $1 + 2 = 3$ | 5.5 | $5.5 < 12 \text{ m}$ | YES |

ece3849_int_latency

Round Robin Polling

- In the Round Robin Polling option the infinite polling loop contains three if statements.
- The Interrupt status of each event is checked to see if an event occurred.
 - It is masked with the TimerA timeout interrupt bit, TIMER_TIMA_TIMEOUT.
 - This separates the event status from other unrelated interrupts.
 - If the Interrupt for that event occurred, the interrupt handler function is called.

```
124 #ifdef ROUND_ROBIN_POLLING
125     while (true) {
126         if (TimerIntStatus(TIMER0_BASE, 1) & TIMER_TIMA_TIMEOUT) { // event 0 has occurred
127             event0_handler();
128         }
129         if (TimerIntStatus(TIMER1_BASE, 1) & TIMER_TIMA_TIMEOUT) { // event 1 has occurred
130             event1_handler();
131         }
132         if (TimerIntStatus(TIMER2_BASE, 1) & TIMER_TIMA_TIMEOUT) { // event 2 has occurred
133             event2_handler();
134         }
135     }
136
137#endif
```

ece3849_int_latency

Example 1: event2 = 2.5

- Run settings

```
33 #define ROUND_ROBIN_POLLING 1
34 // #define PRIORITY_POLLING 1
35
36 // event and handler definitions
37 #define EVENT0_PERIOD 6007 // [us] event0 period
38 #define EVENT0_EXECUTION_TIME 2000 // [us] event0 handler execution time
39
40 #define EVENT1_PERIOD 8101 // [us] event1 period
41 #define EVENT1_EXECUTION_TIME 1000 // [us] event1 handler execution time
42
43 #define EVENT2_PERIOD 12301 // [us] event2 period
44 #define EVENT2_EXECUTION_TIME 2500 // [us] event2 handler execution time
45
46 // build options
47 // #define DISABLE_INTERRUPTS_IN_ISR // if defined, interrupts are disabled in the body of each ISR
48
```

- ## • Run results

Example 2: Event 2 = 3.5 ms

- Lets try Round Robin Again

$$t_{\text{exec0}} = 2 \text{ ms}$$

$$t_{\text{exec1}} = 1 \text{ ms}$$

$$t_{\text{exec2}} = 3.5 \text{ ms}$$

```
void main (void) {
    <init>; // pseudo-code in <...>
    while (1) {
        if (<event0 occurred>) {
            <handle event0>;
        }
        if (<event1 occurred>) {
            <handle event1>;
        }
        if (<event2 occurred>) {
            <handle event2>;
        }
    }
}
```

| Event | Period | Execution Time | Latency (sum of other events execution times) | Response Time (Latency + execution time) | Relative Deadline (Period) | Schedulable ? YES = response < deadline |
|--------|--------|----------------|--|---|-------------------------------|--|
| event0 | 6 ms | 2 ms | $1 + 3.5 \text{ m} = 4.5$ | $4.5 + 2 = 6.5$ | $6.5 > 6$ | NO |
| event1 | 8 ms | 1 ms | $2 + 3.5 = 5.5$ | $1 + 5.5 = 6.5$ | $6.5 < 8$ | YES |
| event2 | 12 ms | 3.5 ms | $2+1 = 3$ | $3.5 + 3 = 6.5$ | $6.5 < 12$ | YES |

ece3849_int_latency

Example 2: event2 = 3.5 ms

- Run settings

```
33 #define ROUND_ROBIN_POLLING 1
34 // #define PRIORITY_POLLING 1
35
36 // event and handler definitions
37 #define EVENT0_PERIOD 6007 // [us] event0 period
38 #define EVENT0_EXECUTION_TIME 2000 // [us] event0 handler execution time
39
40 #define EVENT1_PERIOD 8101 // [us] event1 period
41 #define EVENT1_EXECUTION_TIME 1000 // [us] event1 handler execution time
42
43 #define EVENT2_PERIOD 12301 // [us] event2 period
44 #define EVENT2_EXECUTION_TIME 3500 // [us] event2 handler execution time
```

| (x)= Variables | Expressions | Registers | | | |
|--------------------------------------|--------------|-----------|------------|------------------|--|
| | | | | | |
| Expression | Type | | Value | Expected Results | |
| (x)= event0_latency/120.0f | float | | 4502.19189 | 4.5 ms | |
| (x)= event1_latency/120.0f | float | | 5499.8667 | 5.5 ms | |
| (x)= event2_latency/120.0f | float | | 2998.80835 | 3.0 ms | |
| (x)= event0_response_time/120.0f | float | | 6502.8667 | 6.5 ms | |
| (x)= event1_response_time/120.0f | float | | 6500.5249 | 6.5 ms | |
| (x)= event2_response_time/120.0f | float | | 6499.4834 | 6.5 ms | |
| (x)= event0_missed_deadlines | unsigned int | | 50 | Missed Deadlines | |
| (x)= event1_missed_deadlines | unsigned int | | 0 | Confirmed! | |
| (x)= event2_missed_deadlines | unsigned int | | 0 | | |
| + Add new expression | | | | | |

- Run results

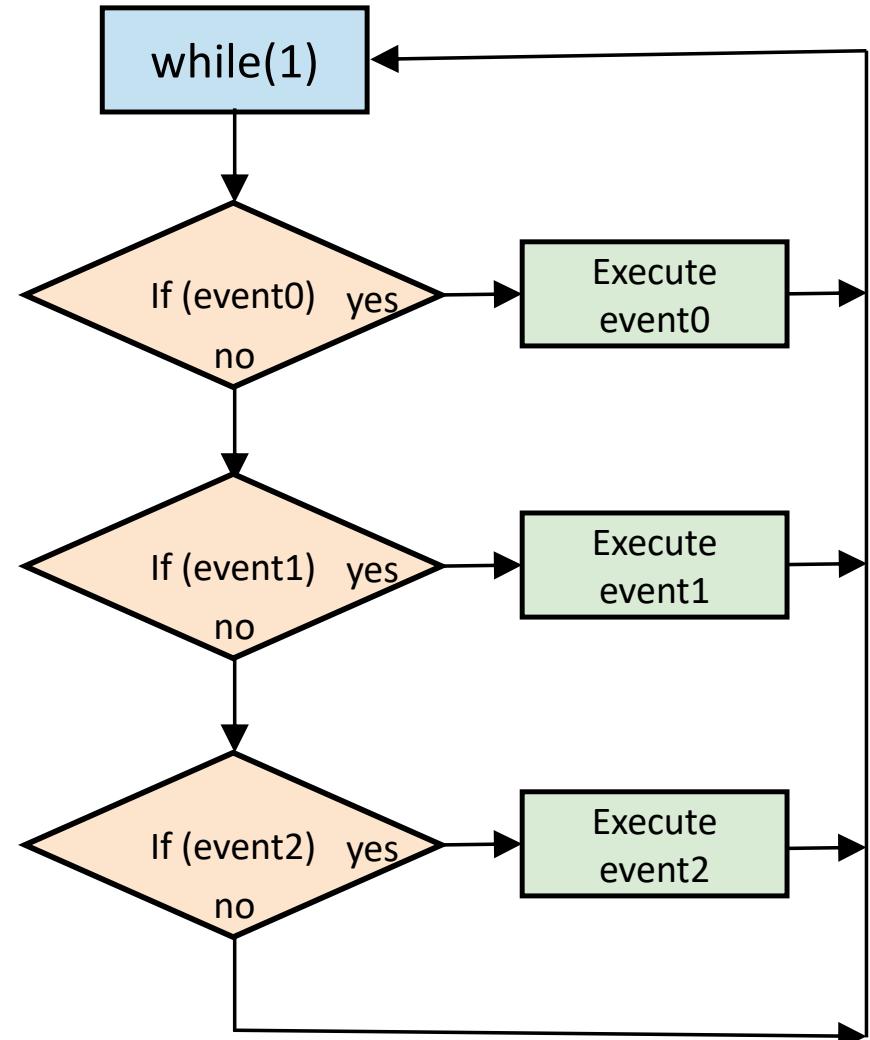
Example 2, event2 = 3.5 ms

Round Robin is not schedulable

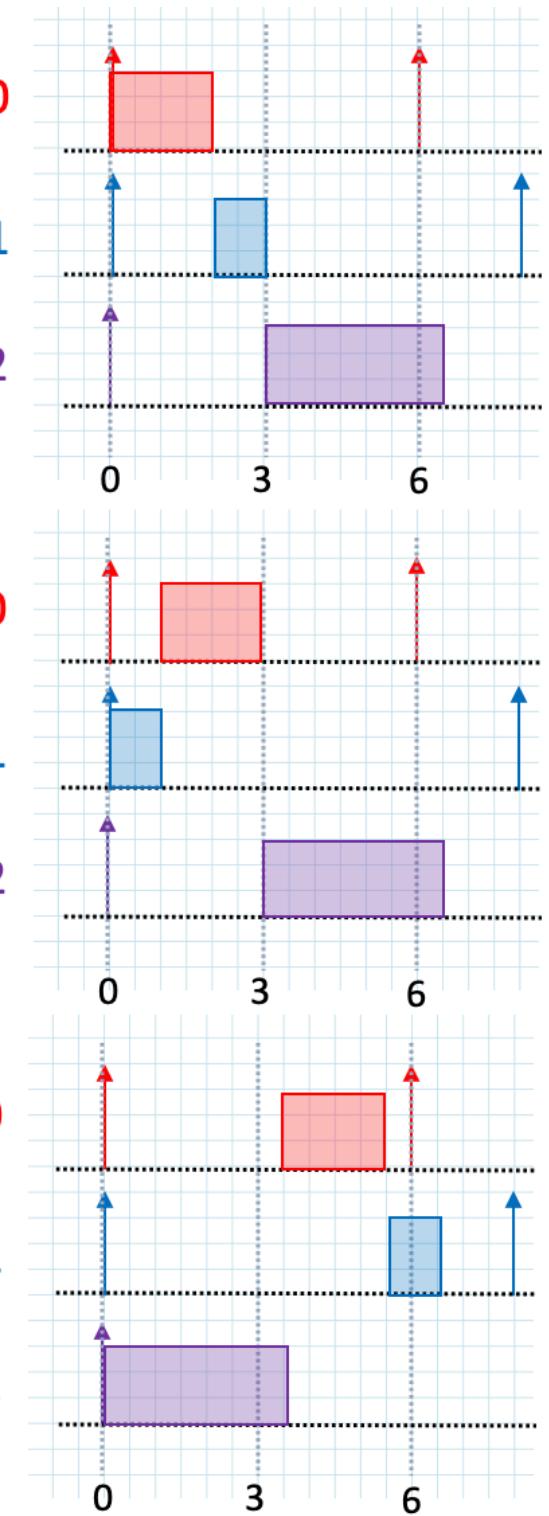
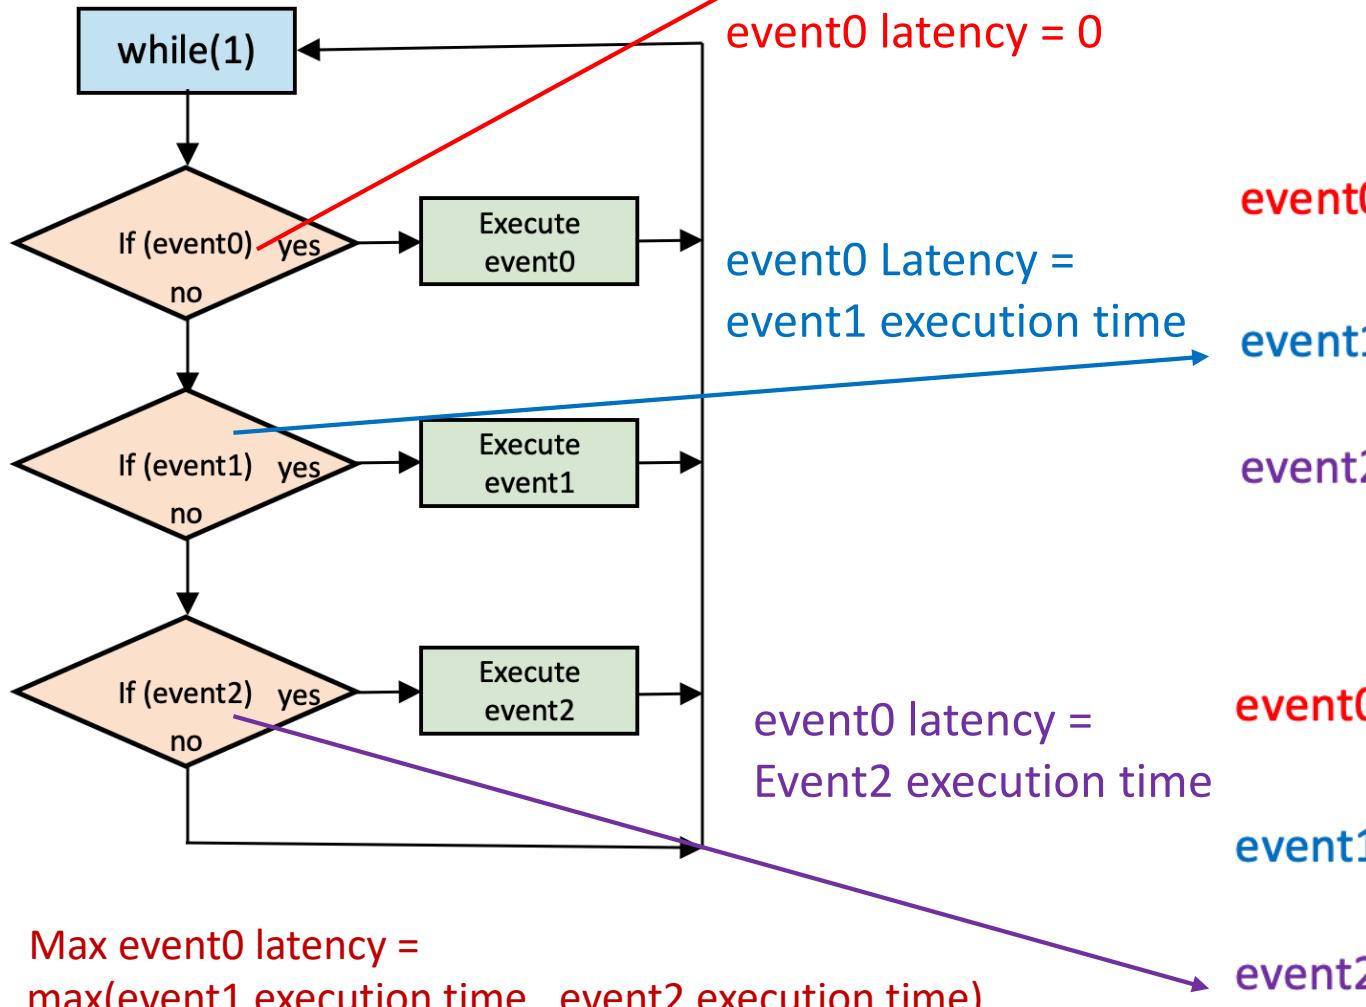
- The execution time of all the events factor into the response time of ALL the events.
- What is the bottleneck in this case?
 - Short event has the most problems?
- How might we schedule differently?
 - Recheck our requirements and see if we could the period long?
 - Different polling – prioritize shortest first.

Priority Polling: Shortest deadline first

- event0 will always be serviced.
- Only if event0 is not waiting will event 1 be serviced.
- Only if event0 and event1 are not waiting will event 2 be serviced.
- What is likely to happen if event0 takes most of the CPU time?
 - Possibly 1 or 2 is starved?



Priority Polling: event0 latency



Priority Polling

```
void main (void) {  
    <init>;  
  
    while (1) {  
        if (<event0 occurred>) {    texec0 = 2 ms  
            <handle event0>;  
        }  
        else if (<event1 occurred>) {  
            <handle event1>;          texec1 = 1 ms  
        }  
        else if (<event2 occurred>) {  
            <handle event2>;          texec2 = 3.5 ms  
        }  
    }  
}
```

| Event | Period | Execution Time | Latency | Response Time (Latency + execution time) | Relative Deadline (Period) | Schedulable ? YES = response < deadline |
|--------|--------|----------------|-------------------------------|---|-------------------------------|--|
| event0 | 6 ms | 2 ms | 3.5 m (max of 1 OR 3.5) | 5.5 | 5.5 < 6 | YES |
| event1 | 8 ms | 1 ms | | | | |
| event2 | 12 ms | 3.5 ms | | | | |

- What will happen if event0 period gets shorter?

- ?

ece3849_int_latency

Example 2: event2 = 3.5, Priority Polling

- Run settings

```
33 // #define ROUND_ROBIN_POLLING 1
34 #define PRIORITY_POLLING 1
35
36 // event and handler definitions
37 #define EVENT0_PERIOD 6007 // [us] event0 period
38 #define EVENT0_EXECUTION_TIME 2000 // [us] event0 handler execution time
39
40 #define EVENT1_PERIOD 8101 // [us] event1 period
41 #define EVENT1_EXECUTION_TIME 1000 // [us] event1 handler execution time
42
43 #define EVENT2_PERIOD 12301 // [us] event2 period
44 #define EVENT2_EXECUTION_TIME 3500 // [us] event2 handler execution time
45
```

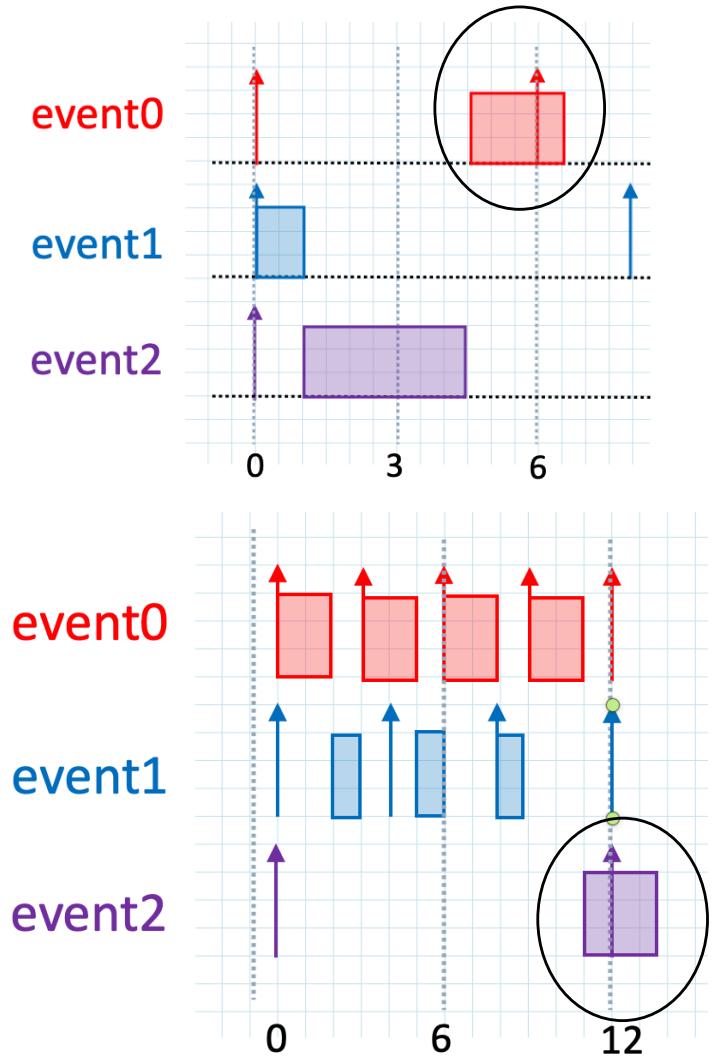
- Run results

| (x)= Variables | Expressions | Registers | | | | | | | | |
|----------------------------------|--------------|-----------|------------|------------------|--|--|--|--|--|--|
| Expression | Type | | Value | Expected Results | | | | | | |
| (x)= event0_latency/120.0f | float | | 3500.92505 | 3.5 ms | | | | | | |
| (x)= event1_latency/120.0f | float | | 5500.2334 | 5.5 ms | | | | | | |
| (x)= event2_latency/120.0f | float | | 3002.07495 | 3.0 ms | | | | | | |
| (x)= event0_response_time/120.0f | float | | 5501.56689 | 5.5 ms | | | | | | |
| (x)= event1_response_time/120.0f | float | | 6500.94189 | 6.5 ms | | | | | | |
| (x)= event2_response_time/120.0f | float | | 6502.75 | 6.5 ms | | | | | | |
| (x)= event0_missed_deadlines | unsigned int | | 0 | | | | | | | |
| (x)= event1_missed_deadlines | unsigned int | | 0 | | | | | | | |
| (x)= event2_missed_deadlines | unsigned int | | 0 | | | | | | | |
| + Add new expression | | | | | | | | | | |

Response time lower
No missed deadlines!

Round Robin vs. Priority Polling

- Round Robin Scheduling
 - Pro: Will guarantee each event gets processed.
 - Con: Response time is limited by the sum of the execution times of all events.
 - May cause events with short deadlines to miss their deadline.
- Prioritize Short Deadline First
 - Pro: Short dead lines will be met.
 - Con: Starvation
 - Events with long deadlines may never get serviced if higher priority tasks occur at a fast enough rate.
 - Figuring out maximum latency for low priority tasks complicated.



Example 3: Event 2 = 6 msec

- Lets try Priority Polling again

$$t_{\text{exec}0} = 2 \text{ ms}$$

$$t_{\text{exec}1} = 1 \text{ ms}$$

$$t_{\text{exec}2} = 6.0 \text{ ms}$$

```
void main (void) {  
    <init>;  
  
    while (1) {  
        if (<event0 occurred>) {  
            <handle event0>;  
        }  
        else if (<event1 occurred>) {  
            <handle event1>;  
        }  
        else if (<event2 occurred>) {  
            <handle event2>;  
        }  
    }  
}
```

| Event | Period | Execution Time | Latency | Response Time (Latency + execution time) | Relative Deadline (Period) | Schedulable ? YES = response < deadline |
|--------|--------|----------------|---------|---|-------------------------------|--|
| event0 | 6 ms | 2 ms | | | | |
| event1 | 8 ms | 1 ms | | | | |
| event2 | 12 ms | 6 ms | | | | |

ece3849_int_latency

Example 3: event2 = 6.0, Priority Polling

- Run settings

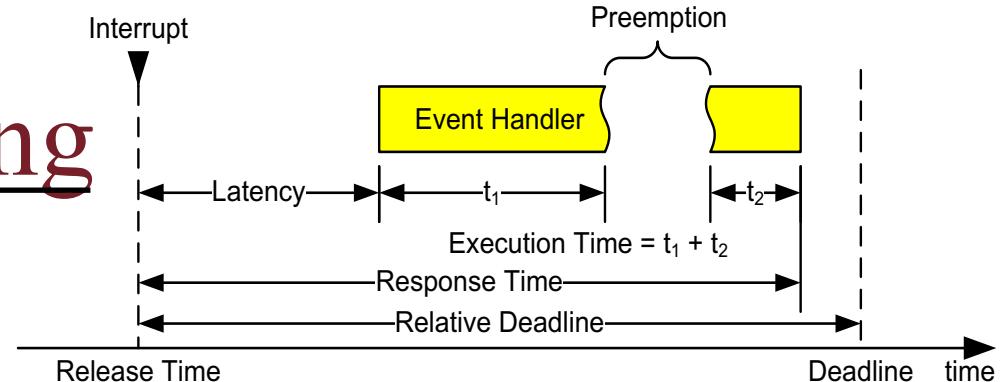
```
33 // #define ROUND_ROBIN_POLLING 1
34 #define PRIORITY_POLLING 1
35
36 // event and handler definitions
37 #define EVENT0_PERIOD 6007 // [us] event0 period
38 #define EVENT0_EXECUTION_TIME 2000 // [us] event0 handler execution time
39
40 #define EVENT1_PERIOD 8101 // [us] event1 period
41 #define EVENT1_EXECUTION_TIME 1000 // [us] event1 handler execution time
42
43 #define EVENT2_PERIOD 12301 // [us] event2 period
44 #define EVENT2_EXECUTION_TIME 6000 // [us] event2 handler execution time
...
```

| (x)= Variables | Expressions | Registers | | | | | | | |
|----------------------------------|--------------|-----------|------------|---------------------------|------------------|--|--|--|--|
| Expression | Type | | Value | | Expected Results | | | | |
| (x)= event0_latency/120.0f | float | | 6000.875 | | 6.0 ms | | | | |
| (x)= event1_latency/120.0f | float | | 8098.4585 | | 8.0 ms | | | | |
| (x)= event2_latency/120.0f | float | | 3002.3833 | | 3.0 ms | | | | |
| (x)= event0_response_time/120.0f | float | | 8001.5415 | | 8.0 ms | | | | |
| (x)= event1_response_time/120.0f | float | | 9099.18359 | | 9.0 ms | | | | |
| (x)= event2_response_time/120.0f | float | | 9003.05859 | | 9.0 ms | | | | |
| (x)= event0_missed_deadlines | unsigned int | | 940 | | | | | | |
| (x)= event1_missed_deadlines | unsigned int | | 332 | Tons of missed deadlines! | | | | | |
| (x)= event2_missed_deadlines | unsigned int | | 0 | | | | | | |
| + Add new expression | | | | | | | | | |

- What strategies can we try next to meet all the deadlines?

- ?

Preemptive scheduling



- The simplest preemptive scheduling requires an interrupt controller.
 - The Cortex-M4 is designed for real time systems and supports an interrupt-driven strategy.
- The infinite while loop runs non-critical tasks in the foreground.
- Interrupt controller automatically calls an **Interrupt Service routine (ISR)** as soon as an event occurs.
 - The foreground tasks are preempted when ISRs are called.
 - ISR's are said to run in the **background**.
 - Use of foreground and background is arbitrary and sometimes swapped.
- Higher priority ISRs can interrupt / preempt lower priority ones.

Interrupt-Driven Design

- Beginning of main

- Initialize Interrupts and hardware.
- While loop (in foreground)

- Performs low priority non-interrupt driven tasks.
- Or it just loops if nothing to do.

- ISR_event0 – highest priority ISR

- ISR_event1 – Mid priority ISR

- ISR_event2 – Lowest priority ISR

Outline of an interrupt-driven real-time system:

```
void main (void) {  
    <init>;  
  
    while (1) {  
        <perform non real-time work>;  
    }  
}  
  
void ISR_event0(void) {  
    <handle event0>;  
}  
  
void ISR_event1(void) {  
    <handle event1>;  
}  
  
void ISR_event2(void) {  
    <handle event2>;  
}
```

- Interrupt controller automatically calls ISR routines to operate in the background

Prioritization of Interrupts

- Higher priority ISR can preempt / interrupt lower priority ISRs.
- How should we prioritize the interrupts?
 - Shorter periods should have higher priority.
 - Smaller deadlines
- Once setup it is possible to turn on and off preemption.
 - When preemption is off, the system behaves like a priority polling design.
 - But we need to respond more quickly than this.

ece3849_int_latency:

Enable Interrupts

- Comment out polling option and enable ISR.

```
33 // #define ROUND_ROBIN_POLLING 1
34 // #define PRIORITY_POLLING 1
35
36 // event and handler definitions
37 #define EVENT0_PERIOD 6007 // [us] event0 period
38 #define EVENT0_EXECUTION_TIME 2000 // [us] event0 handler execution time
39
40 #define EVENT1_PERIOD 8101 // [us] event1 period
41 #define EVENT1_EXECUTION_TIME 1000 // [us] event1 handler execution time
42
43 #define EVENT2_PERIOD 12301 // [us] event2 period
44 #define EVENT2_EXECUTION_TIME 6000 // [us] event2 handler execution time
45
46 // build options
47 #define ENABLE_ISR 1
48 // #define DISABLE_INTERRUPTS_IN_ISR // if defined, interrupts are disabled in t
49
```

ece3849_int_latency:

Enable Interrupts

- ENABLE_ISR allows the IntMasterEnable function to be called.

```
117 #ifdef ENABLE_ISR  
118     IntMasterEnable();  
119 #endif  
....
```

- Creates a while loop with nothing in it.

```
156 #ifdef ENABLE_ISR  
157     while(true) {  
158  
159     }  
160 #endif  
....
```

- In the vector table in tm4c1294ncpdt_startup_css.c the event handler functions are assigned to the Timer A interrupts

```
109     event0_handler,  
110     IntDefaultHandler,  
111     event1_handler,  
112     IntDefaultHandler,  
113     event2_handler,  
....
```

```
.....  
// Timer 0 subtimer A  
// Timer 0 subtimer B  
// Timer 1 subtimer A  
// Timer 1 subtimer B  
// Timer 2 subtimer A  
....
```

ece3849_int_latency:

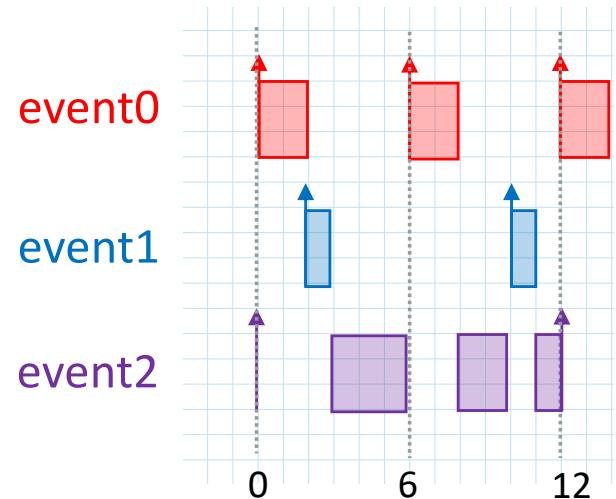
Setting Priority

```
--  
92 // initialize general purpose timers 0-2 for periodic interrupts  
93 SysCtlPeripheralEnable(SYSCTL_PERIPH_TIMER0);  
94 TimerDisable(TIMER0_BASE, TIMER_BOTH);           highest priority = lowest number = 0  
95 TimerConfigure(TIMER0_BASE, TIMER_CFG_PERIODIC);  
96 TimerLoadSet(TIMER0_BASE, TIMER_A, TIMER0_PERIOD - 1);  
97 TimerIntEnable(TIMER0_BASE, TIMER_TIMA_TIMEOUT);  
98 IntPrioritySet(INT_TIMER0A, 0); // 0 = highest priority, 32 = next lower  
99 IntEnable(INT_TIMER0A);  
100  
101 SysCtlPeripheralEnable(SYSCTL_PERIPH_TIMER1);  
102 TimerDisable(TIMER1_BASE, TIMER_BOTH);           Mid priority = 32  
103 TimerConfigure(TIMER1_BASE, TIMER_CFG_PERIODIC);  
104 TimerLoadSet(TIMER1_BASE, TIMER_A, TIMER1_PERIOD - 1);  
105 TimerIntEnable(TIMER1_BASE, TIMER_TIMA_TIMEOUT);  
106 IntPrioritySet(INT_TIMER1A, 32); // 0 = highest priority, 32 = next lower  
107 IntEnable(INT_TIMER1A);  
108  
109 SysCtlPeripheralEnable(SYSCTL_PERIPH_TIMER2);  
110 TimerDisable(TIMER2_BASE, TIMER_BOTH);           Lowest priority = highest number = 64  
111 TimerConfigure(TIMER2_BASE, TIMER_CFG_PERIODIC);  
112 TimerLoadSet(TIMER2_BASE, TIMER_A, TIMER2_PERIOD - 1);  
113 TimerIntEnable(TIMER2_BASE, TIMER_TIMA_TIMEOUT);  
114 IntPrioritySet(INT_TIMER2A, 64); // 0 = highest priority, 32 = next lower  
115 IntEnable(INT_TIMER2A);
```

- The priority value register set by IntPrioritySet function is an 8-bit value.
- However, only the 3 most significant bits of the value are used for interrupt priority.
 - Valid values are 0, 32, 64, 96, 128, 160, 192, 224 (multiples of 32)

Example 3: Event 2 = 6 msec

- Preemptive Interrupt Driven Design
 - Interrupt controller calls ISR immediately on event0.
 - There will be a small delay to call ISR.
 - Higher priority events interrupt/preempt lower priority events.
 - Higher priority events may interrupt multiple times.
 - Events do not have fixed phase and can happen anytime relative to each other.



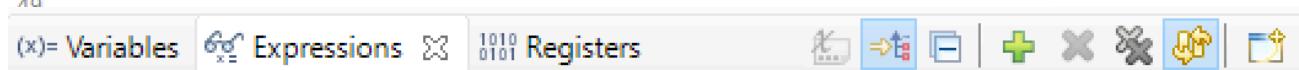
| Event | Period | Execution Time | Latency | Response Time (Latency + ?) | Relative Deadline (Period) | Schedulable ? YES = response < deadline |
|--------|--------|----------------|---------|--------------------------------|-------------------------------|--|
| event0 | 6 ms | 2 ms | | | | |
| event1 | 8 ms | 1 ms | | | | |
| event2 | 12 ms | 6 ms | | | | |

ece3849_int_latency

Example 3: event2 = 6.0, Interrupt Driven

- Run settings

```
33 // #define ROUND_ROBIN_POLLING 1
34 // #define PRIORITY_POLLING 1
35
36 // event and handler definitions
37 #define EVENT0_PERIOD 6007 // [us] event0 period
38 #define EVENT0_EXECUTION_TIME 2000 // [us] event0 handler execution time
39
40 #define EVENT1_PERIOD 8101 // [us] event1 period
41 #define EVENT1_EXECUTION_TIME 1000 // [us] event1 handler execution time
42
43 #define EVENT2_PERIOD 12301 // [us] event2 period
44 #define EVENT2_EXECUTION_TIME 6000 // [us] event2 handler execution time
45
46 // build options
47 #define ENABLE_ISR 1
48 // #define DISABLE_INTERRUPTS_IN_ISR // if defined, interrupts are disabled in t
49
```



| Expression | Type | Value | Expected Results |
|----------------------------------|--------------|-------------|------------------|
| (x)= event0_latency/120.0f | float | 0.266666681 | 0 ms |
| (x)= event1_latency/120.0f | float | 2000.14172 | 2.0 ms |
| (x)= event2_latency/120.0f | float | 3001.25 | 3.0 ms |
| (x)= event0_response_time/120.0f | float | 2000.84998 | 2.0 ms |
| (x)= event1_response_time/120.0f | float | 3001.95825 | 3.0 ms |
| (x)= event2_response_time/120.0f | float | 12005.1504 | 12.0 ms |
| (x)= event0_missed_deadlines | unsigned int | 0 | (marginal) |
| (x)= event1_missed_deadlines | unsigned int | 0 | |
| (x)= event2_missed_deadlines | unsigned int | 0 | |

- Run results

- Why is the event0 latency not 0?

Latency of Highest Interrupt

- If ISR are enabled, the latency of the highest priority interrupt will be small but not zero.
 - The current operation needs to complete.
 - The context of the CPU state is saved.
 - The first instruction of the ISR is fetched.
- Also in this the ece3849_int_latency program several operations were executed before the latency was calculated.
 - The time was read to calculate the latency.
 - This would normally not be considered part of latency.

ece3849_int_latency

Example 3: event2 = 6.0, Interrupt Driven, all the same priority

```
92 // initialize general purpose timers 0-2 for periodic interrupts
93 SysCtlPeripheralEnable(SYSCTL_PERIPH_TIMER0);
94 TimerDisable(TIMER0_BASE, TIMER_BOTH);
95 TimerConfigure(TIMER0_BASE, TIMER_CFG_PERIODIC);
96 TimerLoadSet(TIMER0_BASE, TIMER_A, TIMER0_PERIOD - 1);
97 TimerIntEnable(TIMER0_BASE, TIMER_TIMA_TIMEOUT);
98 IntPrioritySet(INT_TIMER0A, 0); // 0 = highest priority, 32 = next lower
99 IntEnable(INT_TIMER0A);

100
101 SysCtlPeripheralEnable(SYSCTL_PERIPH_TIMER1);
102 TimerDisable(TIMER1_BASE, TIMER_BOTH);
103 TimerConfigure(TIMER1_BASE, TIMER_CFG_PERIODIC);
104 TimerLoadSet(TIMER1_BASE, TIMER_A, TIMER1_PERIOD - 1);
105 TimerIntEnable(TIMER1_BASE, TIMER_TIMA_TIMEOUT);
106 IntPrioritySet(INT_TIMER1A, 0); // 0 = highest priority, 32 = next lower
107 IntEnable(INT_TIMER1A);

108
109 SysCtlPeripheralEnable(SYSCTL_PERIPH_TIMER2);
110 TimerDisable(TIMER2_BASE, TIMER_BOTH);
111 TimerConfigure(TIMER2_BASE, TIMER_CFG_PERIODIC);
112 TimerLoadSet(TIMER2_BASE, TIMER_A, TIMER2_PERIOD - 1);
113 TimerIntEnable(TIMER2_BASE, TIMER_TIMA_TIMEOUT);
114 IntPrioritySet(INT_TIMER2A, 0); // 0 = highest priority, 32 = next lower
115 IntEnable(INT_TIMER2A);
***
```

- What happens if we make them all the same priority?
 - ?

ece3849_int_latency

Example 3: $\text{event2} = 6.0$,

Interrupt Driven, all the same priority

- Interrupts with the same software priority are still given priority in the hardware through the vector table.
- If all events have same software priority, system behaves the same as in priority polling design.

| (x)= Variables | Expressions | Registers | | | | | | | | |
|----------------------------------|--------------|-----------|------------|--|-----------------|---------|--|--|--|--|
| Expression | Type | | Value | | Priority | Polling | | | | |
| (x)= event0_latency/120.0f | float | | 6000.4834 | | 6.0 | ms | | | | |
| (x)= event1_latency/120.0f | float | | 8097.3418 | | 8.0 | ms | | | | |
| (x)= event2_latency/120.0f | float | | 3001.1416 | | 3.0 | ms | | | | |
| (x)= event0_response_time/120.0f | float | | 8001.1001 | | 8.0 | ms | | | | |
| (x)= event1_response_time/120.0f | float | | 9097.94141 | | 9.0 | ms | | | | |
| (x)= event2_response_time/120.0f | float | | 9001.75879 | | 9.0 | ms | | | | |
| (x)= event0_missed_deadlines | unsigned int | | 834 | | Not schedulable | | | | | |
| (x)= event1_missed_deadlines | unsigned int | | 291 | | Not schedulable | | | | | |
| (x)= event2_missed_deadlines | unsigned int | | 0 | | | | | | | |
| + Add new expression | | | | | | | | | | |

Disabled Interrupts in the ISR

- In `ece3849_int_latency`, the interrupts can be disabled in the ISR routines.

```
46 // build options
47 #define ENABLE_ISR 1
48 #define DISABLE_INTERRUPTS_IN_ISR // if defined, interrupts are disabled in the body of each ISR
49

203 void event0_handler(void)
204 {
205 #ifndef DISABLE_INTERRUPTS_IN_ISR
206     IntMasterDisable();
207 #endif
208     uint32_t t = TIMER0_PERIOD - TTMR0_TAR_R; // read Timer A count using direct register
209     if (t > event0_latency) event0_latency = t; // measure latency
210     TimerIntClear(TIMER0_BASE, TIMER_TIMA_TIMEOUT); // clear interrupt flag
211
212     delay_us(EVENT0_EXECUTION_TIME); // handle event0
213
214     if (TimerIntStatus(TIMER0_BASE, 1) & TIMER_TIMA_TIMEOUT) { // next event occurred
215         event0_missed_deadlines++;
216         t = 2 * TIMER0_PERIOD; // timer overflowed since last event
217     }
218     else t = TIMER0_PERIOD;
219     t -= TimerValueGet(TIMER0_BASE, TIMER_A); // read Timer A count using driver
220     if (t > event0_response_time) event0_response_time = t; // measure response time
221 #ifdef DISABLE_INTERRUPTS_IN_ISR
222     IntMasterEnable();
223 #endif
224 }
225
```

Interrupts Globally disabled at start of ISR

Interrupts Globally re-enabled at end of ISR

Effect of Disabled Interrupts in the ISR

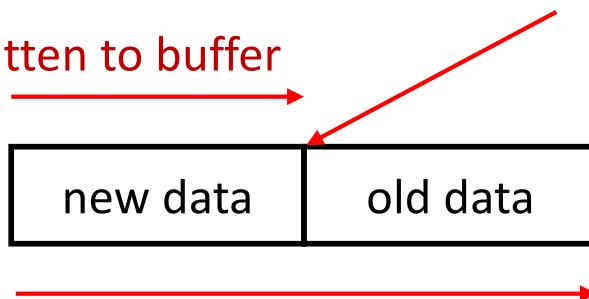
- Under what conditions is it important to disable interrupts?
- What happens to the response time if we do this?
 - ?

Disabling Interrupts / Shared Data

- If ISRs are disabled, the latency will be the longest time they are disabled for.
 - In the ece3849_int_latency program, this was the entire execution time of each ISR.
- We may globally disable interrupts to protect access to shared data and resources.
 - event1 is filling up a buffer.
 - event0 is reading out of the buffer.
 - If event0 preempts event1, then all the data will not get written and the data read out will be a combination of old and new.

#2 Read event preempts the write event.

#1 New data being written to buffer



#3 Read event reads both old and new data causing it to be invalid

ece3849_int_latency

Example 3: event2 = 6.0, Interrupt Driven but disabled in ISR

- If Interrupts are disabled during the entire ISR, no preemption can happen and performance reverts back to priority polling.

| (x)= Variables | Expressions | Registers | | | | | | | | Priority Polling |
|----------------------------------|-------------|--------------|--|--|------------|--|-----------------|--|--|------------------|
| | | | | | | | | | | |
| Expression | | Type | | | Value | | | | | Priority Polling |
| (x)= event0_latency/120.0f | | float | | | 6001.4585 | | | | | 6.0 ms |
| (x)= event1_latency/120.0f | | float | | | 8099.7998 | | | | | 8.0 ms |
| (x)= event2_latency/120.0f | | float | | | 3003.07495 | | | | | 3.0 ms |
| (x)= event0_response_time/120.0f | | float | | | 8002.1499 | | | | | 8.0 ms |
| (x)= event1_response_time/120.0f | | float | | | 9100.5 | | | | | 9.0 ms |
| (x)= event2_response_time/120.0f | | float | | | 9003.72461 | | | | | 9.0 ms |
| (x)= event0_missed_deadlines | | unsigned int | | | 881 | | | | | 0 |
| (x)= event1_missed_deadlines | | unsigned int | | | 312 | | Not schedulable | | | 0 |
| (x)= event2_missed_deadlines | | unsigned int | | | 0 | | | | | 0 |

Disabling Interrupts for short periods

- Maximum latency is affected by disabling interrupts.
 - It is important to minimize the time interrupts are disabled.
 - Only disable them during critical time to avoid shared data problems.
- Example...
 - Disable interrupts only when data being written or read to the buffer.
 - Enable at all other times.
- ece3849 int latency simulation (what do you think will happen?)

```
46 // build options
47 #define ENABLE_ISR 1
48 // #define DISABLE_INTERRUPTS_IN_ISR // if defined, interrupts are disabled in the body of each ISR
49 #define DISABLE_INTERRUPTS_SHORT 1
50
164 #ifdef DISABLE_INTERRUPTS_SHORT
165     //loop for testing the effect of disabling interrupts
166     while (true) {
167         IntMasterDisable();      Turn interrupts off.
168         delay_us(100);        100 usec delay to emulate disabling during data writing.
169         //count_unloaded++;
170         //count_loaded--;
171         IntMasterEnable();    Turn interrupts on.
172
173         IntMasterDisable();    Turn interrupts off.
174         delay_us(200);        200 usec delay to emulate disabling during data reading.
175         IntMasterEnable();    Turn interrupts on.
176     }
177 #endif
```

DISABLE_INTERRUPTS_SHORT Results

| Expression | Type | Value | Interrupt Enable results |
|----------------------------------|--------------|------------|--------------------------|
| (x)= event0_latency/120.0f | float | 200.333328 | 0 ms |
| (x)= event1_latency/120.0f | float | 2149.29175 | 2.0 ms |
| (x)= event2_latency/120.0f | float | 3201.7417 | 3.0 ms |
| (x)= event0_response_time/120.0f | float | 2200.91675 | 2.0 ms |
| (x)= event1_response_time/120.0f | float | 3193.28345 | 3.0 ms |
| (x)= event2_response_time/120.0f | float | 14170.167 | 12.0 ms |
| (x)= event0_missed_deadlines | unsigned int | 0 | (marginal) |
| (x)= event1_missed_deadlines | unsigned int | 0 | |
| (x)= event2_missed_deadlines | unsigned int | 75 | |

Not schedulable
14.2 > 12 msec

- Effect of disabling interrupts for short times.
 - Latency was increased by the maximum disabled time
 - Max(100 usec for write, 200 usec for read) = 200 usec.
 - All Latencies increased by 200 usec.
- Event2 now solidly not schedulable.
 - The latency caused an additional event0 event to preempt event2 adding 2 msec more to its execution time.
 - Designs without margin can be pushed from working to failing miserably by adding even tiny delays.

Even small changes make a big difference.

```
165 //loop for testing the effect of disabling interrupts
166 while (true) {
167     IntMasterDisable();
168 //     delay_us(100);           Removed fixed delays.
169     count_unloaded++;
170     count_loaded--;
171     IntMasterEnable();
172
173     IntMasterDisable();
174 //     delay_us(200);           Now, we are just incrementing and decrement to
175     IntMasterEnable();        variables.
176 }
177 #endif
```

Removed fixed delays.

Now, we are just incrementing and decrement to variables.

We added just 2 lines of code.

- It more than doubled the latency of the highest priority event, 600 nsec >> 266 nsec.

| Expression | Type | New latency | Original Latency |
|----------------------------------|--------------|-------------|------------------|
| (x)= event0_latency/120.0f | float | 0.600000024 | 0.266666681 |
| (x)= event1_latency/120.0f | float | 2000.14172 | 2000.14172 |
| (x)= event2_latency/120.0f | float | 3001.21655 | 3001.25 |
| (x)= event0_response_time/120.0f | float | 2001.20837 | 2000.84998 |
| (x)= event1_response_time/120.0f | float | 3002.2251 | 3001.95825 |
| (x)= event2_response_time/120.0f | float | 12005.3838 | 12005.1504 |
| (x)= event0_missed_deadlines | unsigned int | 0 | 0 |
| (x)= event1_missed_deadlines | unsigned int | 0 | 0 |
| (x)= event2_missed_deadlines | unsigned int | 0 | 0 |

Benchmarks and Metrics

- The latency of the highest priority task is a standard benchmark of a real time system.
 - It has the smallest latency, nothing can do better.
- Definitions of real-time metrics
 - **Interrupt Latency** = From hardware event to running the first instruction of the ISR.
 - **Interrupt Response Time** = Interrupt Latency + Context Saving Time
 - Context saving time is the time to save the contents of the CPU state.
 - **ISR response time** = Interrupt response time + maximum ISR execution time
 - **Interrupt recovery time** = Time to restore the context upon returning from the ISR.

Summary of Scheduling Strategies

- **System assumptions**
 - All events are periodic.
 - The relative deadline = period of the event.
 - Events are not phase aligned.
- **Loops without preemption**
 - The best option to keep things simple.
 - Round-Robin Polling
 - No priority, all events are equal.
 - Every event has an opportunity to execute.
 - Difficult to meet deadline for tasks with small periods.
 - Priority Polling
 - Prioritizes execution of shortest deadline first.
 - Minimizes latency of high priority tasks.
 - If high priority tasks are too frequent or have long execution times, this may cause low priority tasks to starve or miss their deadlines.

Summary of Scheduling Strategies

- Preemptive strategy with interrupts
 - Events with the shortest period are prioritized the highest.
 - They provide the **lowest possible latencies** for the **highest priority** events.
 - Lower priority events still suffer from starvation and are at risk for missing their deadlines.
 - Adds complexity requiring an interrupt controller to trigger events.
 - Higher priority tasks preempt lower priority tasks.
 - Evaluation of the maximum response times becomes difficult.
 - Shared data and resources may be required by multiple ISRs at the same time.
 - This can be addressed by disabling interrupts for very short periods of time.

Looking forward

- We must ensure low enough latency.
 - What we have looked at so far.
 - Interrupts reduce latency but must be prioritized.
 - Minimizing the time interrupts are disabled is critical.
 - Future steps
 - RTOS will provide more advanced scheduling options.
- We must ensure low enough execution time.
 - Profiling code to determine actual performance.
 - Writing code with deterministic execution time.
 - Optimizing code for worst-case performance.
 - Priority inheritance (mutex semaphores).
 - Understanding the hardware that runs the software.
 - Analyzing software performance at the instruction level.
 - Optimizing worst-case memory performance (caching).

Better Analysis Methods

- When we introduced interrupts, response time of lower priority events becomes difficult to calculate.
 - Often they are preempted multiple times?
 - What if conditions change and event 1 can be preempted multiple times too.

| Event | Period | Execution Time | Latency | Response Time (Latency + ?) | Relative Deadline (Period) | Schedulable ? YES = response < deadline |
|--------|--------|----------------|---------|--------------------------------|-------------------------------|--|
| event0 | 6 ms | 2 ms | 0 ms | 2 ms | 6 ms | Yes |
| event1 | 8 ms | 1 ms | 2 ms | 3 ms | 8 ms | Yes |
| event2 | 12 ms | 6 ms | 3 ms | 12 ms Tricky to calc | 12 ms | Maybe, marginal |

- We need a more robust method of calculating response time to guarantee we have thought of the worst case scenario.

Rate-Monotonic Scheduling

Theory

- Rate-Monotonic Scheduling (RMS) Theory provides a more reliable way to verify the scheduling.
- The system must satisfy the following conditions.
 - All real-time tasks (events) are periodic with fixed periods.
 - Relative deadline = event period.
 - Fixed priorities are assigned to events according to period.
 - Shortest period = higher priority.
 - Higher priority tasks preempt lower priority ones.
 - Execution time per event is fixed for each task.
 - This is accomplished by using the maximum value.
 - Task switching overhead is negligible.
 - Tasks do not synchronize with each other
 - No blocking or disabling interrupts.

RMS – Theorem

- Theorem: “The first deadline is the hardest”
- If a task meets its first deadline when all tasks are released at the same time, then it will meet all future ones.
 - Worst case when all the tasks happen at once..
- Look at each task
 - Does it meet its first deadline?
 - If yes, it is schedulable for all conditions.
- Lets use a graphical model.

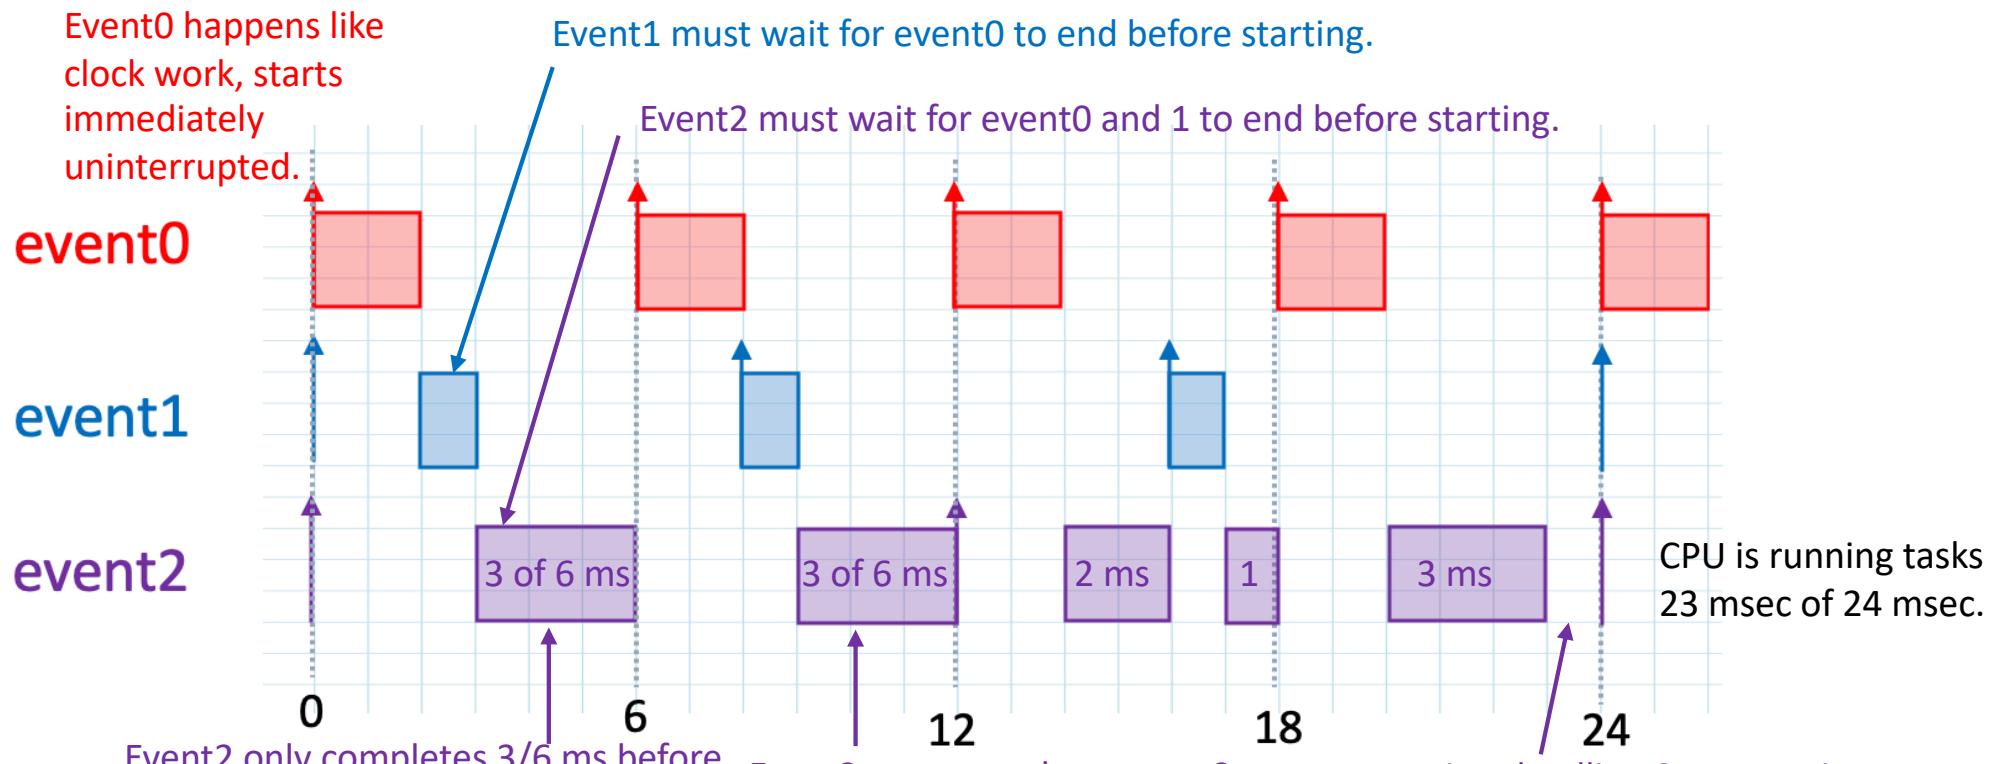
| Task | Period | Execution Time | Priority | Latency | Response Time | Schedulable? |
|-------------|---------------|-----------------------|-----------------|----------------|----------------------|---------------------|
| task0 | 6 ms | 2 ms | | | | |
| task1 | 8 ms | 1 ms | | | | |
| task2 | 12 ms | 6 ms | | | | |

Unrolled schedule:



RMS Graphical Model

| Event | Period | Execution Time | Priority | Latency | Response Time (Latency + ?) | Relative Deadline (Period) | Schedulable ? YES = response < deadline |
|--------|--------|----------------|----------|---------|-----------------------------|----------------------------|--|
| event0 | 6 ms | 2 ms | High | 0 ms | 2 ms | 6 ms | Yes |
| event1 | 8 ms | 1 ms | Mid | 2 ms | 3 ms | 8 ms | Yes |
| event2 | 12 ms | 6 ms | Low | 3 ms | 12 ms | 12 ms | Yes, marginal |



| Task | Period | Execution Time | Priority | Latency | Response Time | Schedulable? |
|-------------|---------------|-----------------------|-----------------|----------------|----------------------|---------------------|
| task0 | 6 ms | 1 ms | | | | |
| task1 | 8 ms | 4 ms | | | | |
| task2 | 12 ms | 3 ms | | | | |

Unrolled schedule:



RMS Graphical Model

| Event | Period | Execution Time | Priority | Latency | Response Time (Latency + ?) | Relative Deadline (Period) | Schedulable ? YES = response < deadline |
|--------|--------|----------------|----------|---------|-----------------------------|----------------------------|--|
| event0 | 6 ms | 1 ms | high | 0 ms | 1 ms | 6 ms | yes |
| event1 | 8 ms | 4 ms | mid | 1 ms | 5 ms | 8 ms | yes |
| event2 | 12 ms | 3 ms | low | 5 ms | 14 ms | 12 ms | NO |

Event0 happens like Event1 must wait for event0 to end before starting.

clock work, starts immediately uninterrupted.

event0

event1

event2

Event2 must wait for event0 and 1 to end before starting.

Event1 is preempted by event0.

CPU is running tasks 22 msec of 24 msec.

Event2 only completes 1/6 ms before event0 and then preempts it.

Event2 resumes when event 0 is complete. Event 1 preempts it.

Event2 resumes when event 0 is complete. Next event2 is waiting to start and is preempted by event1.