

ECE3849  
D-Term 2021

Real Time Embedded Systems

Module 2 Part 1

# Module 2 Part 1 Overview

- Scheduling Strategies Continued.
  - CPU Load Characterization.
  - Earliest Deadline First Scheduling.

# ece3849\_int\_latency: cpu\_load\_count()

- The ENABLE\_CPU\_LOAD\_COUNT option allows the program to estimate the fraction of time spent in real time tasks / ISR running in the background versus non-real time events running in the foreground.
- Program functionality.
  - TIMER3 initialization.
    - Configures a TIMER3 in a one shot configuration with a duration of 1 second.
    - When enabled the timer will count down to 0 over a 1 second interval.
    - Once at count = 0, it sets its interrupt status bit and holds its count at 0.
  - `cpu_load_count()` function.
    - Clears the interrupt status bit.
    - Starts the timer.
    - Then sits in a while loop waiting for the timers interrupt status bit to be set.
    - Each time through the while loop, a loop variable, i, is incremented.

# ece3849\_int\_latency: CPU load calculation

- `cpu_load_count()` returns the number of loops completed in one timer interval = 1 second for this example.
- The number of while loops counted is proportional to the amount of time spent in the foreground.
  - When no interrupts are running,
    - All of the time will be in the foreground and the count will be high.
    - This measurement is taken once at the beginning of the program and stored in the `count_unloaded` variable.
  - When interrupts are enabled,
    - `cpu_load_count` will be called repeatedly in the foreground loop for low priority tasks.
    - When the function is preempted by background ISRs, it does not count.
    - The resulting count is stored in the `count_loaded` variable.
  - The fraction of time spent in the foreground can be calculated by `count_loaded/count_unloaded`.
- CPU Utilization =  $1 - \text{count\_loaded} / \text{count\_unloaded}$ .
- This gives us the average CPU utilization.

# ece3849 int\_latency: CPU load calculation

## main()

```
181 // code for keeping track of CPU load
182 #ifdef ENABLE_CPU_LOAD_COUNT
183
184 // initialize timer 3 in one-shot mode for polled timing
185 SysCtlPeripheralEnable(SYSCTL_PERIPH_TIMER3);
186 TimerDisable(TIMER3_BASE, TIMER_BOTH);
187 TimerConfigure(TIMER3_BASE, TIMER_CFG_ONE_SHOT);
188 TimerLoadSet(TIMER3_BASE, TIMER_A, gSystemClock - 1); // 1
189
190 count_unloaded = cpu_load_count();
191
192 IntMasterEnable();
193
194 TimerEnable(TIMER0_BASE, TIMER_A);
195 TimerEnable(TIMER1_BASE, TIMER_A);
196 TimerEnable(TIMER2_BASE, TIMER_A);
197
198 while (1) {
199     count_loaded = cpu_load_count();
200     cpu_load = 1.0f - (float)count_loaded/count_unloaded;
201 }
```

### Configures TIMER3.

- Enables TIMER3 peripheral for use.
- Disables operation during configuration.
- Configures the timer for ONE\_SHOT mode.
- Loads the count register to 1 second.

Takes one measurement with interrupt.  
disabled initialize count\_unloaded.

Enables Interrupts.

Starts task counters, to  
trigger interrupts.

Enters foreground while loop,  
continually calling cpu\_load\_count and  
calculating CPU utilization.

## cpu\_load\_count

```
282 uint32_t cpu_load_count(void)
283 {
284     uint32_t i = 0;
285     TimerIntClear(TIMER3_BASE, TIMER_TIMA_TIMEOUT);
286     TimerEnable(TIMER3_BASE, TIMER_A); // start one-shot timer
287     while (!(TimerIntStatus(TIMER3_BASE, false) & TIMER_TIMA_TIMEOUT))
288         i++;
289     return i;
290 }
```

- Initializes the loop variable to 0.
- Clears the interrupt status from last time.
- Enables the counter.
- While TIMER3 is counting, it increments i.
- When 1 second is up it returns the i value.

# ece3849\_int\_latency: CPU load calculation

- Run Settings

```
33 // #define ROUND_ROBIN_POLLING 1
34 // #define PRIORITY_POLLING 1
35
36 // event and handler definitions
37 #define EVENT0_PERIOD 6007 // [us] event0 period
38 #define EVENT0_EXECUTION_TIME 2000 // [us] event0 handler execution time
39
40 #define EVENT1_PERIOD 8101 // [us] event1 period
41 #define EVENT1_EXECUTION_TIME 1000 // [us] event1 handler execution time
42
43 #define EVENT2_PERIOD 12301 // [us] event2 period
44 #define EVENT2_EXECUTION_TIME 6000 // [us] event2 handler execution time
45
46 // build options
47 // #define ENABLE_ISR 1
48 // #define DISABLE_INTERRUPTS_IN_ISR // if defined, interrupts are disabled in the
49 // #define DISABLE_INTERRUPTS_SHORT 1
50 #define ENABLE_CPU_LOAD_COUNT 1
51
```

- Run Results

Expression	Type	Value	
(x)= event0_latency/120.0f	float	0.308333337	
(x)= event1_latency/120.0f	float	1998.10828	
(x)= event2_latency/120.0f	float	2988.9917	
(x)= event0_response_time/120.0f	float	2000.83337	
(x)= event1_response_time/120.0f	float	3001.8584	0.94 is greater than the upper bound of 0.7798 for three tasks but still schedulable.
(x)= event2_response_time/120.0f	float	12005.0254	
(x)= event0_missed_deadlines	unsigned int	0	
(x)= event1_missed_deadlines	unsigned int	0	
(x)= event2_missed_deadlines	unsigned int	0	
(x)= cpu_load	float	0.943958461	Expected Value: 0.958

# ece3849\_int\_latency: Missing Deadlines

```
33 // #define ROUND_ROBIN_POLLING 1
34 // #define PRIORITY_POLLING 1
35
36 // event and handler definitions
37 #define EVENT0_PERIOD 6007 // [us] event0 period
38 #define EVENT0_EXECUTION_TIME 1000 // [us] event0 handler execution time
39
40 #define EVENT1_PERIOD 8101 // [us] event1 period
41 #define EVENT1_EXECUTION_TIME 4000 // [us] event1 handler execution time
42
43 #define EVENT2_PERIOD 12301 // [us] event2 period
44 #define EVENT2_EXECUTION_TIME 3000 // [us] event2 handler execution time
45
46 // build options
47 // #define ENABLE_ISR 1
48 // #define DISABLE_INTERRUPTS_IN_ISR // if defined, interrupts are disabled in
49 // #define DISABLE_INTERRUPTS_SHORT 1
50 #define ENABLE_CPU_LOAD_COUNT 1
E1
```

Expression	Type	Value
(x)= event0_latency/120.0f	float	0.308333337
(x)= event1_latency/120.0f	float	998.216675
(x)= event2_latency/120.0f	float	4983.4834
(x)= event0_response_time/120.0f	float	1001.01666
(x)= event1_response_time/120.0f	float	5002.2002
(x)= event2_response_time/120.0f	float	14006.917
(x)= event0_missed_deadlines	unsigned int	0
(x)= event1_missed_deadlines	unsigned int	0
(x)= event2_missed_deadlines	unsigned int	100
(x)= cpu_load	float	0.904288411

We are missing deadlines.

But we are only using 90% of the time?

We have still have 10% CPU time to spare.

# RMS Graphical Model

<b>Task</b>	<b>Period</b>	<b>Execution Time</b>	<b>Priority</b>	<b>Latency</b>	<b>Response Time</b>	<b>Schedulable?</b>
task0	6 ms	1 ms				
task1	8 ms	4 ms				
task2	12 ms	3 ms				

Unrolled schedule:



How might change our scheduling algorithm to meet our deadlines?

- ?

# RMS Graphical Model

Event	Period	Execution Time	Priority	Latency	Response Time (Latency + ?)	Relative Deadline (Period)	Schedulable ? YES = response < deadline
event0	6 ms	1 ms	high	0 ms	1 ms	6 ms	yes
event1	8 ms	4 ms	mid	1 ms	5 ms	8 ms	yes
event2	12 ms	3 ms	low	5 ms	14 ms	12 ms	NO

Event0 happens like Event1 must wait for event0 to end before starting.

clock work, starts immediately uninterrupted.

event0

event1

event2

Event2 must wait for event0 and 1 to end before starting.

Event1 is preempted by event0.

CPU is running tasks 22 msec of 24 msec.

Event2 only completes 1/6 ms before event0 and then 0 preempts it.

Event2 resumes when event 0 is complete. Event 1 preempts it.

Event2 resumes when event 0 is complete. Next event2 is waiting to start and is preempted by event1.

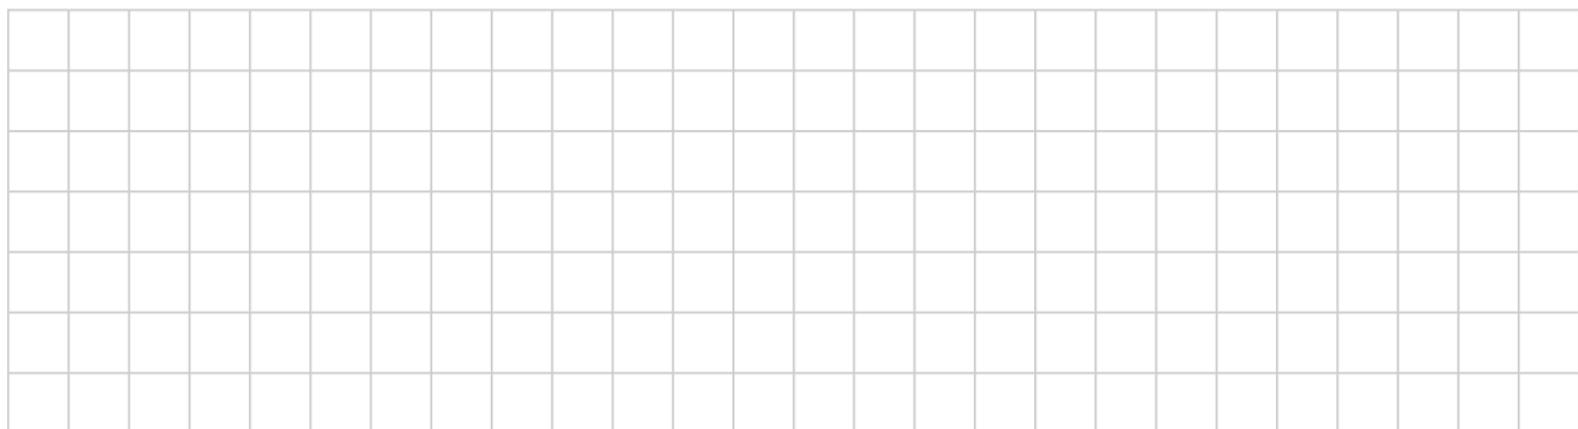
# Earliest Deadline First Scheduling

- For the previous example Period = [6,8,12], Execution time = [1,4,3].
- The CPU utilization is
  - $U = 1/6 + 4/8 + 3/12 = 0.9167 = 91.67\%$
  - We still have extra CPU cycles to use:  $91.67\% < 100\%$
- To meet the deadline we need to throw out the fixed priority condition in the RMS Theory.
- The other conditions are still met
  - All real-time tasks are periodic with fixed periods.
  - Higher priority tasks preempt lower priority ones.
  - Execution time per event is fixed for each task.
  - Task switching overhead is negligible.
  - Tasks do not synchronize with each other.
- Earliest Deadline First (EDF) Scheduling
  - Uses Dynamic Priority
  - Looks at which deadline is coming first and chooses that task to execute.
  - The closer to the deadline, the higher the priority.
  - System needs to know the interrupt periods and keep track of them.

# Earliest Deadline First Scheduling

- Tasks with the same deadline, have the same priority.
  - Tasks with the same priority will not preempt each other for this analysis.
  - In general system there will be a variety of choices on how to handle equal priority interrupts.

Event	Period	Execution Time	Latency	Response Time (Latency + ?)	Schedulable ? YES = response < deadline
event0	6 ms	1 ms			
event1	8 ms	4 ms			
event2	12 ms	3 ms			



# Earliest Deadline First Scheduling

Event	Period	Execution Time	Latency	Response Time (Latency + ?)	Schedulable ? YES = response < deadline
event0	6 ms	1 ms	3 ms	4 ms	Yes
event1	8 ms	4 ms	1 ms	5 ms	Yes
event2	12 ms	3 ms	5 ms	8 ms	Yes

#1 event0 has earliest = 6 , event1 = 8 , event2 = 12

Utilization = 22/24 = 91.67%

#2 event1 has earliest = 8, event0 nothing pending, event2 = 12

#3 event2 has earliest = 12, event0 & 2 nothing pending

#4 event2 and 0 = 12, they are equal complete event2 (no preemption)

#5 event0 has earliest = 12 , event1 = 16 , event2 nothing pending

#6 event1 has earliest = 16, event0 & 2 nothing pending

#7 event0 has earliest = 18 , event2 = 24, event1 nothing pending

#8 event2 has earliest = 24, nothing pending on 0 & 1

#9 event2 & 1 has earliest = 24, equal no preemption.

#10 event1 has earliest = 24, nothing pending on 0 or 2

#11 event1 & 0 = 24, equal no preemption.

#12 event0 has earliest = 24 ,12 nothing pending on 1 or 2

event0

event1

event2

0

6

12

18

24

# Looking forward

- We must ensure low enough latency.
  - What we have looked at so far.
    - Interrupts reduce latency but must be prioritized.
    - Minimizing the time interrupts are disabled is critical.
  - Future steps
    - RTOS will provide more advanced scheduling options.
- We must ensure low enough execution time.
  - Profiling code to determine actual performance.
  - Writing code with deterministic execution time.
  - Optimizing code for worst-case performance.
  - Priority inheritance (mutex semaphores).
  - Understanding the hardware that runs the software.
  - Analyzing software performance at the instruction level.
  - Optimizing worst-case memory performance (caching).