

ECE3849
D-Term 2021

Real Time Embedded Systems

Module 2 Part 2

Module 2 Part 2 Overview

- Handling Shared Data
 - Disabling Interrupts
 - Priority Inversion
 - Atomic Operations
 - Using Local Variables
 - Common Shared Data Bugs

Shared Data and Resources

- The only way to pass data to or from an ISR is through global variables.
 - This is referred to as inter-thread communication.
 - These globals are called shared data.
 - Data can be shared between ISRs or between the foreground and background tasks.
 - Issues arise when multiple threads try to access the same resource, particularly if one is reading and the other writing.
- Peripherals are classified as shared resources, and treated in the same way as shared data.

Shared Data Problem: Time Example

```
int iHours, iMinutes, iSeconds;
```

```
void TimerISR(void) {  
    // clear interrupt flag  
    if (++iSeconds >= 60) {  
        iSeconds = 0;  
        if (++iMinutes >= 60) {  
            iMinutes = 0;  
            if (++iHours >= 24)  
                iHours = 0;  
        }  
    }  
}
```

```
long SecondsSinceMidnight(void) {  
    return ((iHours * 60) + iMinutes) * 60 + iSeconds;  
}
```



Global variables = shared data



Timer ISR called once per second and increments the hours, minutes and seconds.

- Time sequence -hh:mm:ss

03:59:58

03:59:59

04:00:00

04:00:01

Converts time back to seconds.

- Function called whenever needed.

Do we have any control over when the function is called relative to the ISR?

- ?

Any concerns?

- ?

Shared Data Problem: Time Example

```
int iHours, iMinutes, iSeconds; →  
  
void TimerISR(void) {  
    // clear interrupt flag  
    if (++iSeconds >= 60) {  
        iSeconds = 0;  
        if (++iMinutes >= 60) {  
            iMinutes = 0;  
            if (++iHours >= 24)  
                iHours = 0;  
        }  
    }  
}  
  
long SecondsSinceMidnight(void) {  
    return ((iHours * 60) + iMinutes) * 60 + iSeconds;  
}
```

Timer ISR called once per second and increments the hours, minutes and seconds.

- Time sequence -hh:mm:ss

03:59:58

03:59:59

04:00:00

04:00:01

What if it is 3:59:59 and the interrupt occurs here?

iHours = ?

iMinutes = ?

iSeconds = ?

Do we have a bug: ?

If so, how often will it happen?

Shared Data Fix: Disable Interrupts

```
#include "inc/hw_types.h"
#include "driverlib/interrupt.h"

long SecondsSinceMidnight(void) {
    long lTemp;

    IntMasterDisable();
    lTemp = ((iHours * 60) + iMinutes) * 60 + iSeconds; // Critical section.
    IntMasterEnable();
    return lTemp;
}
```

Disable the Interrupt so ISR can not change values.

Critical section.

Re-enable interrupts after critical section.

Interrupts will resume execution.

ISR will be on a waitlist if they happen during this time.

- It is possible to disable the interrupts during the critical calculation.
 - Care must be taken to minimize the time interrupts are disabled.
 - Why?
- What happens if interrupts are originally disabled?
 - ?

Shared Data : Disable Interrupts Fix

```
#include "inc/hw_types.h"
#include "driverlib/interrupt.h"

long SecondsSinceMidnight(void) {
    tBoolean bIntAlreadyDisabled;
    long lTemp;

    bIntAlreadyDisabled = IntMasterDisable();
    lTemp = ((iHours * 60) + iMinutes) * 60 + iSeconds;
    if (!bIntAlreadyDisabled)
        IntMasterEnable();
    return lTemp;
}
```

IntMasterDisable returns a Boolean that contains the original state.

TRUE: if originally disabled.
FALSE: if originally enabled

Now only enable interrupts if they were originally enabled.

- **IntMasterDisable() and IntMasterEnable() are specific to the TI libraries.**
 - If you ever wanted to move the code to a different processor. Every instance of them would need to be renamed everywhere in the code.
 - Portability is the ease of which code is moved from one environment to another.
 - Adding the macros below will improve portability as you only need to change one line of code instead of every instance.

```
#define FunctionDisablesInterrupts() tBoolean bIntAlreadyDisabled;
#define EnterCritical() bIntAlreadyDisabled = IntMasterDisable();
#define ExitCritical() if (!bIntAlreadyDisabled) IntMasterEnable();
```

Disabling Interrupts Pros & Cons


- Pros
 - Easy to solve the shared data problem.
 - Just a couple lines of code
 - Doesn't take long to execute, CPU load not affected to much
- Cons
 - Additional latency is introduced.
 - The longest period of time that interrupts are disabled is added to the latency of every ISR.
 - If other parts of the system have longer disabled periods, the maximum latency will not be increased.

Disabling a Specific Interrupt

```
long SecondsSinceMidnight(void) {  
    long lTemp;  
  
    IntDisable(INT_TIMER0A);  
    lTemp = ((iHours * 60) + iMinutes) * 60 + iSeconds;  
    IntEnable(INT_TIMER0A);  
    return lTemp;  
}
```

Disable just the interrupt we care about.
Example TIMER0A is used to increment
Time.

Small increase in latency of high priority task0.




- Expected behavior
 - Task 0 and 2 share TIMER0A.
 - Task2 disables TIMER0A Interrupt.
 - Task0 latency increases by a the small delay of the critical section.

- Task1 does not share TIMER0A. How will it affect latency?
 - ?

Priority Inversion

- Task0 and Task2 share TIMEROA Interrupt.
 - Task0 will be blocked by Task2 during critical section.
- Task1 does not share TIMEROA
 - Task1 can preempt Task2 during the critical section.



• Pros of Disabling just TIMEROA

- Less impact in overall system performance.
- For example, if we are sampling data with an ADC in task1, it is unaffected by the TIMEROA interrupt.

• Cons

- Increases complexity making it more error prone.
- Opens the possibility of priority inversion.

Atomic Operations

- What if we simplify the functionality to just deal with seconds?

- We increment seconds in the ISR and roll over at midnight to 0.
- Then simply return one value, lSeconds, in the function call.
- Use a different function to calculate hours, minutes, seconds when needed.

```
long lSeconds;  
  
void TimerISR(void) {  
    // clear interrupt flag  
    if (++lSeconds >= 60 * 60 * 24) {  
        lSeconds = 0;  
    }  
}  
  
long SecondsSinceMidnight(void) {  
    return lSeconds;  
}
```

- lSeconds is a 32-bit value, we have a 32-bit CPU.
 - lSeconds can be read in one clock cycle.
 - Operations that can be completed in a single clock cycle are called **atomic operations**.
 - There is no shared data problem in this case.
- What if lSeconds was a 64-bit value would we have a problem?
 - ?.

TimerISR: 1Seconds Simplification

```
long lSeconds;

void TimerISR(void) {
    // clear interrupt flag
    if (++lSeconds >= 60 * 60 * 24) {
        lSeconds = 0;
    }
}

long SecondsSinceMidnight(void) {
    return lSeconds;
}
```

- What if SecondsSinceMidnight is called from a higher priority interrupt than TimerISR?
 - ?

- This fix depends both on the data type and the hardware architecture.

- This makes it hard to port to other CPUs,
- Easy to make a mistake.
- Difficult to detect and debug.

Effects of Interrupting TimerISR

```
long lSeconds;  
  
void TimerISR(void) {  
    // clear interrupt flag  
    if (++lSeconds >= 60 * 60 * 24) {  
        lSeconds = 0;  
    }  
}  
  
long SecondsSinceMidnight(void) {  
    return lSeconds;  
}
```

- **++lSeconds**
 - Reads lSeconds
 - Modifies lSeconds by incrementing.
 - Writes new value of lSeconds.
- Executes a second write function, to clear the lSeconds value.

- After incrementing, lSeconds == 60*60*24
 - This is not a valid value, not expected in the final result.
 - If the high priority interrupt occurs between the two writes, this value will be returned instead of the expected value of 0.

Local Variable Fix

- Re-writing the TimerISR as shown below
 - Reduces the number of writes to the global variable to only one atomic write.
 - No shared data problem.

```
if (lSeconds + 1 >= 60*60*24)
    lSeconds = 0;           → lSeconds + 1 is stored in a temporary local
else
    lSeconds++;            → lSeconds global variable only written once.
```

Common Shared Data Bugs

	Low-priority thread = main()	high-priority thread = ISR
Operation	Low-priority thread	High-priority thread
Read-modify-write	<ol style="list-style-type: none"> 1. Reads the global variable producing a local copy 2. Modifies the local copy 3. Writes the modified copy to the global variable 	Writes to the global variable

If high priority write occurs between steps 1 and 3, data is lost.

Non-atomic read or multiple reads	<ol style="list-style-type: none"> 1. Reads part of the global variable 2. Reads another part of the global variable 	Writes to the global variable
-----------------------------------	--	-------------------------------

If high priority write occurs between steps 1 and 2.

$\frac{1}{2}$ of the read value is old and $\frac{1}{2}$ of the read value is new.

Non-atomic write or multiple writes	<ol style="list-style-type: none"> 1. Writes part of the global variable 2. Writes another part of the global variable 	Reads the global variable
-------------------------------------	--	---------------------------

If high priority read occurs between steps 1 and 2.

$\frac{1}{2}$ of the read value is old and $\frac{1}{2}$ of the read value is new.

Write followed by read	<ol style="list-style-type: none"> 1. Writes to the global variable 2. Reads the global variable, expecting the originally written value 	Writes to the global variable
------------------------	--	-------------------------------

If high priority write occurs between steps 1 and 2.

The read value is not the expected value written by the lower priority thread.