

ECE3849
D-Term 2021

Real Time Embedded Systems

Module 2 Part 3

Module 2 Part 3 Overview

- Handling Shared Data
 - Atomic Operations
 - Avoid multiple read / write
 - Cortex-M4 Bit Banding functions
 - Inter-thread communications
 - Binary Semaphore
 - Mailboxes
 - Buffering

Atomic Operations

- Operations that can be completed in a single instruction are called atomic operations.
- Global variables that are larger than a single 32-bit CPU register are not atomic in the Cortex-M4.
- For other processors, this may be different.
 - You need to know the number of bits in your primitive data type.
 - You need to know CPU architecture.
 - How many bits is the CPU?
 - Does it have special instructions that might allow normally non-atomic operations to be atomic?
 - Boolean and char are always reads and writes are always atomic independent of hardware architecture.
- Using type definitions helps
 - Using type definitions like `uint32_t` makes code clearer and more portable.
 - Using `int`, `short`, `long` which may vary in length between 16, 32, and 64-bit processors.

Atomic Operations

- Are these Atomic or Non-Atomic?
 - `x++`
 - ?
 - `x += y;`
 - ?
 - `x = y;`
 - ?
 - `if (x > y) { x = 0; }`
 - ?
 - `x |= 1 << 3;`
 - ?
- Not sure check the disassembled code.

Read-Modify-Writes

- Watch out for implicit read-modify-writes
 - $x += y;$ $\rightarrow x = x + y;$ (reads x, adds y , writes x)
 - $x--;$ $\rightarrow x = x + 1;$ (reads x, adds 1, writes x)
 - Single line of code but multiple instructions.
 - If an interrupt occurs after the modify but before the write. The ISR gets over written and is lost.
- Some read-modify-writes may not have shared data problems.
 - main() (low priority task) is doing a read-modify-write.
 - ISR can read it once using an atomic read and store it in a local variable for further use in the ISR.

Cortex-M4F Bit-banding

- Bit-banding allows atomic writes to a single bit in a 32-bit integer.
- In Cortex-M4F, the HWREGBITW() MACRO performs bit-banding.
 - Non-Atomic: $x |= 1 << 3;$
 - $x = x | 0x8;$ //0x8 = binary 1000
 - Atomic equivalent:
 - `uint32_t x;`
 - `HWREGBITW(&x, 3) = 1;` // sets third bit location
 - Not portable, specific to this hardware
 - To use: `#include "inc/hw_types.h"`

Read and Write Recommendations

- **Never write to the same global twice in a row.**
 - If a ISR reads between the writes, it may grab an intermediate value.
- **Avoid reading the same global variable more than once if another thread writes it.**

```
y = gValue + 5;
```

```
...
```

```
z = gValue + 8;
```

Interrupt occurs writes new value
z will not equal y + 3 as expected.

- **If interrupts are infrequent, it may be possible to perform a non-atomic read.**
 - This is accomplished by reading the desired value repeatedly until the same value is returned twice.
 - If the same value is read twice, you have verified an ISR did not occur during the operation.

Repeated Non-Atomic Reading

- Make the variable **volatile** to let the compiler know it might be changed at any time.

```
volatile int iHours, iMinutes, iSeconds; // shared variables
```

```
void TimerISR(void) {  
    // clear interrupt flag  
    if (++iSeconds >= 60) {  
        iSeconds = 0;  
        if (++iMinutes >= 60) {  
            iMinutes = 0;  
            if (++iHours >= 24)  
                iHours = 0;  
        }  
    }  
}
```

TimerISR is only called
once per second.
(very slow rate)

Declare two sets of variables
to hold the values from two
sets of read operations.

```
long SecondsSinceMidnight(void) { // called from low-priority thread  
    int iHr, iMin, iSec; // local copies of the shared data  
    int iHr2, iMin2, iSec2; // second copy
```

If ISR occurred during these
statements an incorrect
value would be returned.

- Read them **until** two
consecutive reads match.
- This confirms no ISR happened.

```
iHr = iHours; iMin = iMinutes; iSec = iSeconds; // read globals  
do {  
    iHr2 = iHr; iMin2 = iMin; iSec2 = iSec;  
    iHr = iHours; iMin = iMinutes; iSec = iSeconds;  
} while (iHr2 != iHr || iMin2 != iMin || iSec2 != iSec);
```

```
return ((iHr * 60) + iMin) * 60 + iSec; // use local copies
```

Calculates return value with confirmed good data.

Repeated Non-Atomic Reading

- What would happen if we left **volatile** out?

- ?

- Why don't we get stuck in the loop?

- ?

- For complex data structures, this can be inefficient and error prone.

- A sequence number can be added to the data.
 - In the ISR
 - The sequence number is incremented each time it is called
 - In the function
 - The sequence number is read.
 - The data is read.
 - The sequence number is read a second time and compared to the original.
 - If the sequence number is different, the data is re-read until the sequence number remains the same.

```
volatile int iHours, iMinutes, iSeconds; // shared variables

void TimerISR(void) {
    // clear interrupt flag
    if (++iSeconds >= 60) {
        iSeconds = 0;
        if (++iMinutes >= 60) {
            iMinutes = 0;
            if (++iHours >= 24)
                iHours = 0;
        }
    }
}

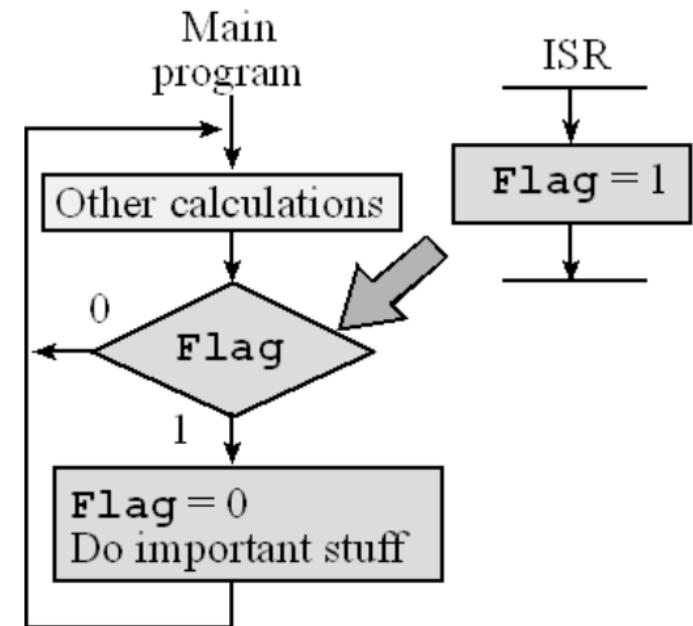
long SecondsSinceMidnight(void) { // called from low-priority thread
    int iHr, iMin, iSec;      // local copies of the shared data
    int iHr2, iMin2, iSec2; // second copy

    iHr = iHours; iMin = iMinutes; iSec = iSeconds; // read globals
    do {
        iHr2 = iHr; iMin2 = iMin; iSec2 = iSec;
        iHr = iHours; iMin = iMinutes; iSec = iSeconds;
    } while (iHr2 != iHr || iMin2 != iMin || iSec2 != iSec);

    return ((iHr * 60) + iMin) * 60 + iSec; // use local copies
}
```

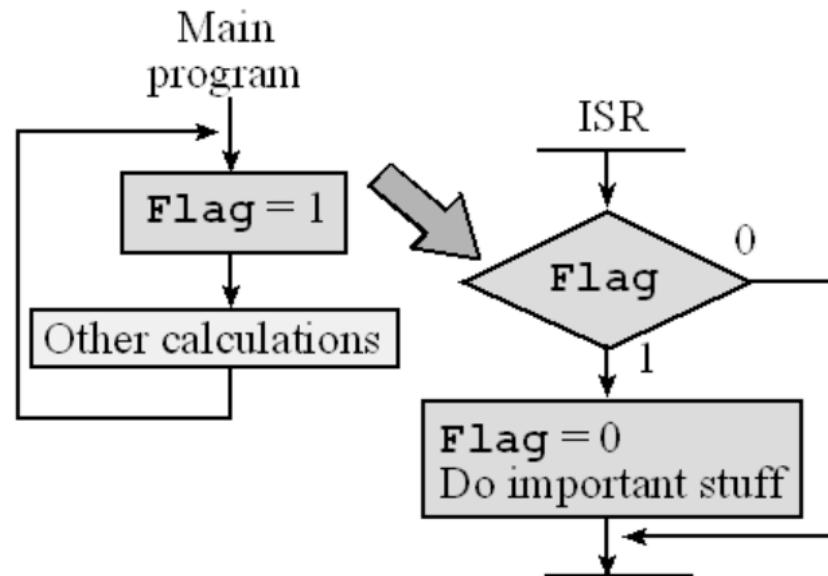
Binary Semaphore

- A **binary semaphore** protects access to shared resources using a shared flag.
 - One thread performs the **signal** action
 - Sets the flag.
 - One thread performs the **wait** action,
 - Waits for the flag to be set.
 - Clears the flag.
 - Performs shared data operations.
- The signal action can be in the ISR and the wait action in the lower priority thread.
- Example:
 - ISR is reading an ADC value.
 - ISR sets flag when new data is available.
 - main() is calculating the average.
 - Waits for the flag.
 - Reads the new ADC value into a local variable
 - Clears the flag.
 - Calculates the new average.



Binary Semaphore

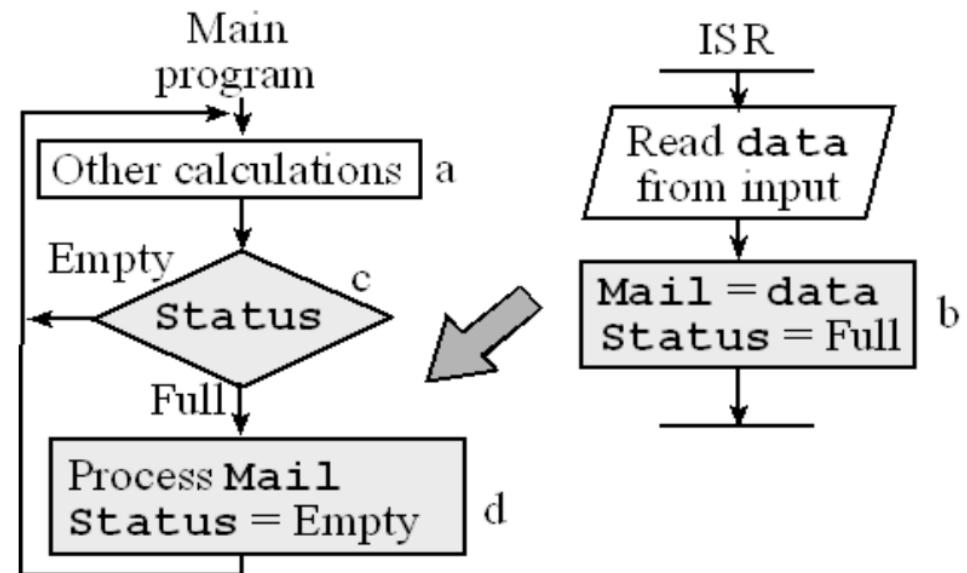
- The signal action can be in the main function or lower priority ISR and the wait action in the real time ISR.



- Example:
 - ISR is responsible for reading a voltage monitor status.
 - main() detects an error in the status and sets the flag.
 - On the next ISR
 - The ISR clears the flag.
 - Reads the voltage status like it does every time but also reads the voltage itself to get more information on the failure.

Mailboxes

- Mailboxes are binary semaphores, that protect a specific data variable. They have two shared global variables.
 - Mail which contains the shared data.
 - Status which is the binary semaphore flag.
- The Mail can be a single variable, multiple variables, structures, etc. Whatever is required in the application.
- Sequence of events.
 - (a) main() is busy doing low priority functions.
 - (b) ISR functionality
 - Reads data from input.
 - Places data in Mail.
 - Sets Status = Full.
 - main()
 - (c) Polls Status.
 - If it is Full, reads the data from Mail into local variable.
 - (d) Set Status = Empty.



Mailbox: Lab0 Example

- Lab 0 has a shared data problem.
 - gButtons is read multiple times during the `snprint` statement for the LCD.
 - If a button press interrupt occurs while writing the button status string in `main()`. The wrong value could be displayed.
- Solution: Mailbox
 - Button ISR functions.
 - Waits for button press.
 - Button info is placed in `gButtons` (global variable).
 - Sets `gMailboxStatus = 1.` (global variable status = Full)
 - Main function polls status, if it is Full,
 - Reads `gButtons` into local variable.
 - Sets `gMailboxStatus = 0.` (global variable status = Empty)
 - Prints string using local variable.

Mailbox: Code Example

Define global Mail and Status variables

ISR performs signal action

main() performs wait action

Example 2: Simple mailbox

```
int bMailboxFull = 0; ←  
// arbitrary size mailbox data structure  
  
void SomeISR(void) {  
    // clear interrupt flag  
    if (!bMailboxFull) { ←  
        // put data into mailbox  
        bMailboxFull = 1;  
    }  
}  
  
void main(void) {  
    // init  
    while (1) {  
        if (bMailboxFull) { ←  
            // extract data from mailbox  
            bMailboxFull = 0;  
        }  
    }  
}
```

Initializes Status to Empty

- If Empty, writes data to Mail and sets flag.
- If Full, no action taken.

- If Full, reads data and clears flag.
- If Empty, no action taken.

- What timing conditions between threads must be met for data not to be lost?
 - ?

Buffering

- Data can be buffered to avoid reading and writing to the same location.
 - Data is written to sequential locations in an array by one thread.
 - When the end of the array is reached, the index wraps around to index 0.
 - Data is read from the buffer by another thread.
 - The timing of the reads and writes are controlled to guarantee there is no overlapping access.
- Buffering (lab 1, ADC sampling routines)
 - High and low priority threads access different parts of the buffer.
 - ADC writing into one half of the buffer.
 - The oscilloscope function reads the buffer index and chooses to process the data in the part of the buffer not being written to.