

ECE3849
D-Term 2021

Real Time Embedded Systems

Module 2 Part 4

Module 2 Part 4 Overview

- Handling Shared Data
 - Circular Buffers: FIFOs

Circular Buffer: FIFO


- FIFO stands for First In First Out.
 - FIFOs are implemented with a global data array with fixed maximum length.
 - Two threads share the FIFO, one reading and one writing.
 - As long as the reading rate can keep up with the average writing rate, no data will be lost if the FIFO is sized properly.
- Several globals variables are required to implement a FIFO.

```
#define FIFO_SIZE 11          // FIFO capacity is 1 item fewer
typedef char DataType;       // FIFO data type
DataType fifo[FIFO_SIZE];   // FIFO storage array
volatile int fifo_head = 0;  // index of the first item in the FIFO
volatile int fifo_tail = 0;  // index one step past the last item
```

- We write the FIFO with a function `tito_put()`, which increments the `fifo_tail` index.
- We read the FIFO with a function `fifo_get()`, which increments the `fifo_head`.

FIFO: Put and Get operations

Initialized: FIFO Empty
`fifo_head = fifo_tail = 0.`




fifo_put(): Writes into FIFO.


- Increments new_tail value.
- If `fifo_head != new_tail`
 - writes data, increments `fifo_tail` & is not full

fifo_get(): Reads out of FIFO.

- If `fifo_head != fifo_tail`,
 - It reads data.
 - Increments the head value.
 - Returns not empty.
- If reads again in this state, the FIFO returns an empty status.



Say you write 9 more times, the tail would be at 10.
If the head or the tail is at the maximum index, the index rolls back to zero .



FIFO Full State: `fifo_tail + 1 == head`
If a `fifo_put()` is issued now no further data can be written. The tail is not incremented and full is returned.

FIFO: Code Example

```
// put data into the FIFO, skip if full
// returns 1 on success, 0 if FIFO was full
int fifo_put(DataType data)
{
    int new_tail = fifo_tail + 1;
    if (new_tail >= FIFO_SIZE) new_tail = 0; // wrap around
    if (fifo_head != new_tail) {
        fifo[fifo_tail] = data;           // store data into the FIFO
        fifo_tail = new_tail;            // advance FIFO tail index
        return 1;                      // success
    }
    return 0;      // full
}

// get data from the FIFO
// returns 1 on success, 0 if FIFO was empty
int fifo_get(DataType *data)
{
    if (fifo_head != fifo_tail) {
        *data = fifo[fifo_head];        // read data from the FIFO
        fifo_head++;                  // advance FIFO head index
        if (fifo_head >= FIFO_SIZE) fifo_head = 0; // wrap around
        return 1;                      // success
    }
    return 0;      // empty
}
```


- Local variable new_tail contains incremented tail value.
If the new_tail is at its maximum ivalue, it wraps around to zero.
- If the new_tail == head then FIFO is full and nothing is written.
- If it is not full then writes data and updates the fifo_tail to the new_tail value.

- If fifo_head == fifo_tail then the FIFO is empty and Return 0 meaning no data was read
If it is not empty then it...
 - reads data value.
 - Increments the fifo_head value.
 - If the incremented head value is at its maximum, then head wraps around to zero.
- It returns 1 meaning data was read.

- Do we have any shared data problems?
 - ?

Shared data: fifo_put()

```
// put data into the FIFO, skip if full
// returns 1 on success, 0 if FIFO was full
int fifo_put(DataType data)
{
    int new_tail = fifo_tail + 1;
    if (new_tail >= FIFO_SIZE) new_tail = 0;
    if (fifo_head != new_tail) {
        fifo[fifo_tail] = data;           // store
        fifo_tail = new_tail;            // advance
        return 1;                        // success
    }
    return 0;                          // full
}
```



If fifo_get() interrupt, it will see old tail value and will not be aware that the fifo_put() is in process.

fifo_head is read only once & atomic.
fifo_tail is written only once & atomic.

- Do we have any shared data problems?
 - No problems.
 - FIFO may report empty while fifo_put() is in process.
 - Once completed fifo_get() will need to be called again to pickup the new data.

Shared data: fifo_get()


```
// RETURNS 1 ON SUCCESS, 0 IF FIFO WAS FULL
int fifo_put(DataType data)
{
    int new_tail = fifo_tail + 1;
    if (new_tail >= FIFO_SIZE) new_tail = 0; // wrap around
    if (fifo_head != new_tail) { // if the FIFO
        fifo[fifo_tail] = data; // store data in
        fifo_tail = new_tail; // advance FIFO
        return 1;
    }
    return 0; // full
}
```

No problem: If fifo_put() happens here, tail will either think it is full or will increment fifo_tail and add data. fifo_get() still has data to read.

```
int fifo_get(DataType *data)
{
    if (fifo_head != fifo_tail) { // if the FIFO
        *data = fifo[fifo_head]; // read data from
        fifo_head++; // advance FIFO
        if (fifo_head >= FIFO_SIZE) fifo_head = 0;
        return 1;
    }
    return 0; // empty
}
```

If interrupt has two consecutive fifo_put() calls
new_tail should = fifo_head = FULL.

BUG: head incorrect, new_tail writes to location 0 and now thinks it is not full and can overwrite the data in the whole buffer.



Shared data: get_fifo(): fix

```
// get data from the FIFO
// returns 1 on success, 0 if FIFO was empty
int fifo_get(DataType *data)
{
    if (fifo_head != fifo_tail) { // if the FIFO
        *data = fifo[fifo_head]; // read data fr
        fifo_head++; // advance FIFO
        if (fifo_head >= FIFO_SIZE) fifo_head = 0;
        return 1;
    }
    return 0; // empty
}
```

fifo_head +1 stored in local variable.


```
if (fifo_head >= FIFO_SIZE-1) {
    fifo_head = 0;
}
else {
    fifo_head++;
}
```

Writes fifo_head only once in an atomic operation,
no intermediate values possible.

- In addition, you could read fifo_head into a local variable, local_head, similar to what is done in the fifo_put() function.
 - Although not required reduces the number of accesses to the global variables and the potential for shared data bugs.

Common FIFO Usage

- FIFO are commonly used to communicate between high priority (low latency) ISRs and low priority (high latency) ISRs or main() functions.
- Often one task can be bursty, delivering a large packet of data quickly and then nothing for a while. While the other task processes data at a regular interval.
 - FIFOs are perfect for **rate matching** these high priority burst accesses with the slow and steady performance of a lower priority task.
 - The larger the bursts or the longer the latency, the larger the FIFO needs to be.



Example: ece3849_shared_data

- Similar to ece3849_int_latency, the program calculates latency, response_time and missed deadlines.
- It uses the fifo_put() and fifo_get() functions from the lecture slides.
- It has only one ISR: event0_handler().
 - The ISR uses fifo_put() to write up to 5 characters at a time in alphabetical order to the FIFO.
 - The ISR is configured to trigger every 10 msec.
- In the main() while loop for low priority tasks.
 - The loop waits 100msec.
 - Calls fifo_get() once to read a single character from the FIFO.
 - If there is a character in the FIFO,
 - It prints the character to the LCD screen.
 - Increments the location of the string for the next character.
 - If the location is at the last character on the screen it clears the screen and starts printing at the top.

ece3849 _ shared _ data ISR code

Create a static variable to hold next character value.

- Loops 5 times
- If there is room in the FIFO, `fifo_put()` writes the character into the FIFO
 - Increments the character value for next time.
 - If the character is U wraps back around to A.

```
133 void event0_handler(void)
134 {
135     uint32_t t;
136     t = TIMER0_PERIOD - TIMER0_TAR_R; // read Timer A count using direct register access
137     if (t > event0_latency) event0_latency = t; // measure latency
138     TimerIntClear(TIMER0_BASE, TIMER_TIMA_TIMEOUT); // clear interrupt flag
139
140     int i;
141     static char c = 'A';
142     for (i = 0; i < 5; i++) { // attempt to put multiple items into the FIFO
143         if (fifo_put(c)) {
144             c++; // go to the next char only if put was successful
145             if (c > 'U') c = 'A';
146         }
147     }
148     delay_us(EVENT0_EXECUTION_TIME); // handle event0
149
150     if (TimerIntStatus(TIMER0_BASE, 1) & TIMER_TIMA_TIMEOUT) { // next event occurred
151         event0_missed_deadlines++;
152         t = 2 * TIMER0_PERIOD; // timer overflowed since last event
153     }
154     else t = TIMER0_PERIOD;
155     t -= TimerValueGet(TIMER0_BASE, TIMER_A); // read Timer A count using driver
156     if (t > event0_response_time) event0_response_time = t; // measure response time
157 }
```

Event0_handler writes A through U in alphabetical order up to 5 characters if the FIFO is not full.

Static Variables

- **Static variables**
 - They behave similar to global variable except their scope is limited to the function or block they are declared in.
 - They stay in memory from call to call and preserve their value. They are not dynamically allocated.
 - Static variables are stored in the data segment instead of the stack.
- **Lab 1:**
 - It is important to minimize what is on the stack because the stack size is extremely limited.
 - Stack overflows will result in a Fault ISR and demonstrate in unpredictable ways.
 - Large data structures should be declared either globally or using static so as not to run out of stack space.

ece3849 shared data: main()

Waits 100 msec.


- Reads the FIFO.
 - If the FIFO is not empty, prints the character to the LCD buffer.
 - Increments the LCD screen location for the next character.
 - If it is at the bottom of the screen. Clears the screen and sets the next character location to the top left of the LCD.
 - Updates the LCD display with the buffered data.
-
- ```
112 // main loop
113 while (1) {
114 delay_us(100000); // 0.1 sec
115 if (fifo_get(&c)) {
116 GrStringDraw(&sContext, &c, /*length*/ 1, /*x*/ x, /*y*/ y, /*opaque*/ false);
117 x += 6;
118 if (x > LCD_HORIZONTAL_MAX - 6) {
119 x = 0;
120 y += 8;
121 if (y > LCD_VERTICAL_MAX - 8) {
122 y = 0;
123 GrContextForegroundSet(&sContext, ClrBlack);
124 GrRectFill(&sContext, &rectFullScreen); // fill screen with black
125 GrContextForegroundSet(&sContext, ClrYellow);
126 }
127 }
128 }
129 GrFlush(&sContext); // flush the frame buffer to the LCD
130 }
131 }
```

# Functional results

- Default conditions with shared data bug in fifo\_get().
  - No errors seen, A through U printed normally.
  - Chance of hitting bug extremely low.
- We increase chances of hitting error by adding a delay, just after the fifo\_head++ statement.

```
62 int fifo_get(DataType *data)
63 {
64 if (fifo_head != fifo_tail) { // if the FIFO is not empty
65 *data = fifo[fifo_head]; // read data from the FIFO
66 // IntMasterDisable();
67 fifo_head++; // advance FIFO head index
68 | delay_us(1000); // circled here
69 if (fifo_head >= FIFO_SIZE) fifo_head = 0; // wrap around
70 // IntMasterEnable();
71 return 1; // success
72 }
73 return 0; // empty
```

- With the delay we can see 11 missing characters when the bug occurs.
- To fix the issue, we have two options.
  - Disable interrupts around critical area
  - Implement the fix in get\_fifo() from lecture.



# Summary of Shared Data Handling

- Preemptive scheduling using interrupts and ISRs can give high performance / low latency at the cost of shared data problems.
- There are several ways to handle shared data problems.
  - Disabling interrupts guarantees mutually exclusive accesses, only one task at a time can access the data.
    - Pro: Easy to implement.
    - Con: Increases latency.
    - Con: Disabling specific interrupts can lead to priority inversion.
  - Careful Coding and Analysis
    - Use local variables to remove shared data possibilities.
    - Avoid the most common shared data bugs by using atomic operations and minimizing multiple reads and writes to global variables.
    - Con: Difficult to verify proper implementation, bugs can be very infrequent and intermittent.
  - Synchronization techniques
    - Con: Adds complexity and protocol, can be resource intensive requiring additional variables and increasing response times.
    - Con: Need to understand rate at which tasks are happening to guarantee no data loss.
    - Reading multiple times/sequence numbers: Used when interrupts are infrequent, to verify an interrupt did not occur and change the data in the middle of the critical data section.
    - Buffering – limit reading from one part of the buffer while writing to another part.
    - Binary Semaphores – using a Boolean flag to control when shared data accesses and critical functions occur.
    - Mailboxes – using a binary Semaphore to limit access to specific data and data structures.
    - FIFOs – allows read and write data tasks to operate independently at different times and rates without conflicts.