

ECE3849
D-Term 2021

Real Time Embedded Systems

Module 3 Part 3

Module 3 Part 3 Overview

- Deadlock.
- Task Priority Inversion.
- Gates for mutual exclusion.
- Semaphores for signaling and synchronization.

Deadlock

- Deadlock occurs when tasks are pending on each other and thus neither are able to post. They are forever blocking each other.

Let `sem_a` and `sem_b` be two semaphores initialized to count = 1 before the tasks

`sem_a` guards the global variable `a`
`sem_b` guards the global variable `b`

```
int a;  
int b;
```

```
0 void task1(UArg arg0, UArg arg1) { // high priority  
    // init  
    while (1) {  
        ...  
        4 Semaphore_pend(sem_a, BIOS_WAIT_FOREVER); 3  
        Semaphore_pend(sem_b, BIOS_WAIT_FOREVER);  
        a = b;  
        Semaphore_post(sem_b);  
        Semaphore_post(sem_a);  
        ...  
    }  
  
    void task2(UArg arg0, UArg arg1) { // low priority  
        // init  
        while (1) {  
            ...  
            5 Semaphore_pend(sem_b, BIOS_WAIT_FOREVER); 1  
            Semaphore_pend(sem_a, BIOS_WAIT_FOREVER); 2  
            b = a;  
            Semaphore_post(sem_a);  
            Semaphore_post(sem_b);  
            ...  
        }  
    }
```

ISR unblocks task 1

Both tasks are blocked, waiting for the other to post.

0 System initialization.

- `sem_a & sem_b`, count = 1.
- task1 is blocked waiting on ISR.

1 task2 starts and pends on `sem_b`, countB = 1 so it can run.

- countB decremented, countB = 0.
- Task2 continues to run.

2 ISR unblocks task1 (high priority).

- task2 is placed in ready state.
- task1 starts running.

3 task1 pends on `sem_a`, countA = 1 so it can run.

- countA decremented, countA = 0.
- task1 continues.

4 task1 pends on `sem_b`, countB = 0.

- task1 is blocked.
- task2 starts running.

5 task2 pends on `sem_a`, countA = 0.

- task2 is blocked.


Ece3849_rtos_latency: Deadlock

- Ece3849_rtos_latency – deadlock simulation mode.
 - #define ENABLE_DEADLOCK 1
 - Adds the code from the deadlock on previous page, to event0_func and event2_func.
 - Event1_func is left unchanged
 - We observe the event counters to verify which events are still running.
 - Event1 continues to run for the whole test.
 - Event0 and 2 stop incrementing almost immediate, showing they are deadlocked.
- If you are using two semaphores they should always pend in the same order in both tasks.
 - We change Event 2 so that both events pend in the same order

```
Semaphore_pend(sem_a, BIOS_WAIT_FOREVER);  
Semaphore_pend(sem_b, BIOS_WAIT_FOREVER);
```
 - Operation returns to normal and all events counter increment for entire test.


Task Priority Inversion

- Task priority inversion.
 - Low priority task, task_low, takes the semaphore0, protecting the shared data.
 - High priority task, task_high, preempts task_low, runs and pends on the semaphore0 and is blocked.
 - Medium priority task, task_mid, preempts task_low, runs and completes.
 - Low priority task continues, posting to semaphore0 to finally unblock task_high.
- Response time of task_high increases by execution time of task_mid.
- Scheduling priorities are violated and relative deadlines maybe missed.



Priority Inheritance

- Priority inheritance.
 - Temporarily raises the priority of the task holding the semaphore.
 - The priority is raised to the highest priority task blocked on the semaphore.
 - Once the high priority task is unblocked, priorities return to normal.
- Restrictions.
 - Priority inheritance is restricted to use with binary semaphores with the role of protecting critical sections.
- Priority Inheritance prevents task_mid from running while task_high is blocked.



Priority Inheritance: GateMutexPri()

- Ti-RTOS implements priority inheritance using a specialized object **GateMutexPri()**.
 - **GateMutexPri_enter()** is placed at the start of the critical section.
 - **GateMutexPri_leave()** is placed at the end of the critical section.
 - These are used in place of **Semaphore_pend()** and **Semaphore_post()** calls.
- **GateMutexPri_enter()** returns a key
 - The key is later used by **GateMutexPri_leave()** to restore the thread preemption to the state just prior to entering.

```
IArg keyMutexPri0;  
...  
keyMutexPri0 = GateMutexPri_enter(mutexPri0);  
// critical section  
GateMutexPri_leave(mutexPri0, keyMutexPri0);
```

Ece3849_rtos_latency: GateMutexPri()

- ece3849_rtos_latency: adds the following statements to event0 and event2 to simulate protecting critical section.

```
130 #ifdef PRIORITY_INVERSION
131     Semaphore_pend(sem_a, BIOS_WAIT_FOREVER);
132     delay_us(100); // "critical section"
133     Semaphore_post(sem_a);
134 #endif

135 #ifdef ENABLE_GATEMUTEXPRI
136     key = GateMutexPri_enter(gateMutexPri0);
137     delay_us(100); // "critical section"
138     GateMutexPri_leave(gateMutexPri0, key);
139 #endif
```

Critical section protected by Semaphore

OR


Critical section protected by GateMutexPri

Products > SYSBIOS > Synchronization > Gates > GateMutexPri - Module Settings

GateMutexPri is a mutex gate (it can only be held by one thread at a time) which implements priority inheritance in order to prevent priority inversion. Priority inversion occurs when a high priority task has its priority effectively 'inverted' because it is waiting on a gate held by a low priority task.

Add the GateMutexPri module to my configuration

Icon will give full description of functions



BIOS
Clock
Config (ti.drivers)
Config (ti.mw)
Defaults
Error
GateMutexPri
gateMutexPri0

Ece3849_rtos_latency: Highest Priority Task Response Time

Expected task0 execution time results:


- No critical section (original conditions):
 - Response time = 18 usec latency + 2 msec execution time = ~2 msec.
- Critical section protected with Semaphore sem_a => causes priority inversion.
 - Expected response time = ~3.2 msec.
 - + 18 usec original latency +
 - + 2 msec task0 original execution time
 - + 1 msec task 1 execution time
 - + 100 usec task2 critical data section protection
 - + 100 usec task0 critical data section execution.
- Critical section protected with GateMutexPri() => removes priority inversion.
 - Expected response time = ~2.1 msec
 - Semaphore response time of ~3.2 – 1.1 msec task1 execution time because priority inversion is removed.

GateMutex Object

- The GateMutex object is GateMutexPri without the priority inheritance function.
 - It is a binary semaphore for mutual exclusion.
 - It enables or disables access to shared resources.
 - They can not be used by Hwi or Swi threads.
 - It uses `GateMutex_enter()` and `GateMutex_leave()` functions to implement this functionality.
- It behaves similar to a semaphore except, it can be entered multiple times in the same function without blocking.
 - If a semaphore pends multiple times, a post must occur in between each pend to unblock it.
 - If `GateMutex_enter()` is called multiple times in the same task and it already holds the key, it continues without being blocked.
 - It is only used for **mutual exclusion** and can not be used for task synchronization.
 - **Mutual exclusion** guarantees only one thread at a time can access the shared resource.

Priority Inheritance Limitations

- Priority Inheritance is on a per Gate basis.
- When multiple gates are used in a single task they are not aware of each others inherited traits.
 - This can lead to further priority inversion.
- task_high is preempted by task_mid below.



GateTask: Disabling Task Switching

- Task switching can be disabled.
 - Disabling switching will increase task latency.
 - Latency will increase by the longest time tasks switching is disabled.
 - This is similar to the effect that disabling interrupts has on interrupt latency.
 - Uses fewer CPU cycles.
 - No priority inversion or deadlock.
- The GateTask object disables task switching.
 - GateTask_enter() function disables the task switching
 - Returns a key to be used in restoring the thread preemption state when leaving the GateTask.
 - GateTask_leave() function enables task switching and restores task state to the thread preemption state just prior to entering.
- These functions can not be called in Hwi or Swi interrupts.

```
IArg keyGateTask0;  
...  
keyGateTask0 = GateTask_enter(gateTask0);  
// critical section  
GateTask_leave(gateTask0, keyGateTask0);  
...
```

Ece3849_rtos_latency: GateTask

- event0_func and event2_func now call the GateTask functions instead of the Semaphore or GateMutexPri functions.

```
#ifdef ENABLE_GATETASK
    key = GateTask_enter(gateTask0);
    delay_us(100); // "critical section"
    GateTask_leave(gateTask0, key);
#endif
```

Latency now increased from 18 usec to 108 usec.

Latency = Original latency + maximum time switching is disabled.

(x)= event0_latency/120.0f	float	108.300003
(x)= event1_latency/120.0f	float	2230.6001
(x)= event2_latency/120.0f	float	3152.80835
(x)= event0_response_time/120.0f	float	2227.3584
(x)= event1_response_time/120.0f	float	3240.16675
(x)= event2_response_time/120.0f	float	14573.333

Response time task0 similar to GateMutexPri functionality.

RTOS: Gating Interrupts

- Ti-RTOS can also disable interrupts that it controls.
- GateSwi object.
 - Uses `GateSwi_enter()` and `GateSwi_leave()` functions and a key to restore the state on leaving.
 - They can be called from Swi or tasks.
 - GateSWi object disables Swi and task scheduling.
 - Affects latency of all Swi and tasks.
- GateHwi object.
 - Uses `GateHwi_enter()` and `GateHwi_leave()` functions and a key to restore the state on leaving.
 - They can be called from Task, Swi or Hwi threads.
 - Guarantees exclusive access to the CPU.
 - **Does not affect “zero-latency interrupts”.**
 - Affects latency of all Hwi, Swi and tasks.

Summary of Objects used for Mutual Exclusion

Mutual Exclusion: only one thread can access a shared resource at a time.

Object	Pros and Cons
1. Semaphore (pend & post) <ul style="list-style-type: none">• not optimal for mutual exclusion• has other applications	+ targeted - slow (higher CPU load) - does not work with ISRs - priority inversion - deadlock
2. GateMutex (enter & leave) <ul style="list-style-type: none">• binary semaphore for mutual exclusion• cannot be used for signaling/synchronization	+ can be “entered” multiple times by the same task without blocking otherwise same as Semaphore
3. GateMutexPri <ul style="list-style-type: none">• binary semaphore with priority inheritance• tasks are unblocked in priority order• preferred for longer critical sections	+ fixes most priority inversion cases otherwise same as GateMutex
4. GateTask <ul style="list-style-type: none">• disables Task scheduling• preferred for short, frequently occurring critical sections	+ faster (lower CPU load) + no priority inversion + no deadlock - affects latency of all Tasks (but not Swi or Hwi)
5. GateSwi <ul style="list-style-type: none">• disables Swi & Task scheduling	+ works with Swi and Tasks - affects latency of all Tasks and Swi
6. GateHwi <ul style="list-style-type: none">• disables Hwi, Swi & Task scheduling• does not disable “zero-latency interrupts”	+ works with Hwi, Swi and Tasks - affects latency of all Hwi, Swi and Tasks

Semaphores for Signaling

- Semaphores allow Hwi, Swi and other tasks to signal a task.
 - A task is unblocked when an event occurs.
- Below Hwi issues a `Semaphore_post()` to unblock Task1 which is pending.
 - If the Hwi is periodic this allows the task to run at a regular interval.
 - This keeps the CPU load under 100% because pending on semaphores do not take CPU cycles.
- Choice of semaphore type.
 - If this is a binary semaphore, then Hwi interrupts will be missed if the task gets unexpectedly delayed and the interrupt runs multiple time between Task1 pends.
 - A counting semaphore ensures that task1 is called once for every interrupt occurrence even if intermittently it misses its relative deadline.

Initialize semaphore `sem0` to count = 0 before it is first used.

```
void Hwi1(UArg arg0) { // hardware interrupt under TI-RTOS
    // communicate with the hardware
    Semaphore_post(sem0);
}

void Task1(UArg arg0, UArg arg1) {
    // init
    while (1) {
        Semaphore_pend(sem0, BIOS_WAIT_FOREVER);
        // handle event
    }
}
```

Semaphore Signaling: Multiple Tasks

- During Initialization,

- semaphore count = 0.
- Task1 runs first because it is higher priority and pends. It is blocked and put on the wait list.
- Task2 runs second it pends and is placed on the wait list.

```
void Hwi1(UArg arg) { // hardware interrupt under TI-RTOS
    // communicate with the hardware
    Semaphore_post(sem0);
    Semaphore_post(sem0);
}
void Task1(UArg arg0, UArg arg1) { // high priority
    // init
    while (1) {
        Semaphore_pend(sem0, BIOS_WAIT_FOREVER);
        // other code
    }
}
void Task2(UArg arg0, UArg arg1) { // low priority
    // init
    while (1) {
        Semaphore_pend(sem0, BIOS_WAIT_FOREVER);
        // other code
    }
}
```

Some time later Hwi1 interrupt occurs. It posts twice.

- The first post unblocks Task1.
 - It is higher priority AND first in the list.
 - Task1 is placed in Ready state.
- The second post unblocks Task2 and it is placed in Ready state.
- When Hwi1 exits, the scheduler runs.
 - Task1 is higher priority and enters Running state.
 - Task1 completes its loop and pends again and is blocked.
 - Task2 enters running state, completes its loop and pends.

Semaphore Signaling: Multiple Tasks

- What happens if Hwi1 only posts once?

```
void Hwi1(UArg arg) { // hardware interrupt under TI-RTOS
    // communicate with the hardware
    Semaphore_post(sem0);
    Semaphore_post(sem0);
}
void Task1(UArg arg0, UArg arg1) { // high priority
    // init
    while (1) {
        Semaphore_pend(sem0, BIOS_WAIT_FOREVER);
        // other code
    }
}
void Task2(UArg arg0, UArg arg1) { // low priority
    // init
    while (1) {
        Semaphore_pend(sem0, BIOS_WAIT_FOREVER);
        // other code
    }
}
```

First Hwi occurs...


- ?

Second Hwi occurs...

- ?

Semaphore for synchronization

- **Rendezvous:** We would like two tasks to arrive at a specific point in their code before either can continue.



Steps 0 -3: Just once on initialization

- 0 System initialization
 - Sem1 & sem2, count = 0
 - task1 is higher priority
- 1 task1 starts and posts to sem1.
 - task2 has not run yet, nothing on waitlist
 - sem1 count is incremented.
- 2 task1 pends on sem2.
 - sem2 count = 0, task1 is put on wait list.
 - task1 is blocked.
- 3 task2 runs and posts to sem2.
 - task1 is in waitlist and moved to Ready state.
- 4 task1 is higher priority and runs its while loop and posts to sem1.
- 5 task1 pends on on sem2 and is blocked.
- 6 Task2 starts running,
 - takes sem1 and decrements sem1 count.
 - Continues its while loop and then posts to sem2.